

Climate Risk Assessment and Management

Consolidated Lecture Notes

James Doss-Gollin

Table of contents

Welcome	3
Motivation and Scope	3
History	3
Aim	3
How to Use This Resource	4
Structure	4
Prerequisites	4
 I About this book	 5
License	6
Contributing	7
Citing	8
Further Reading	9
Inspiration	9
Stats + ML basics	9
Applications	9
More Stats + ML	10
 Preface	 11
What is climate risk?	11
Risk management	11
Exposure and vulnerability	12
Climate hazard	12
What are good strategies?	12
The simple story	12
Why this isn't enough	13
The stakes of getting it wrong	13
This book	14
What this book is not	14
 II I: Foundations	 16
1 Climate science	17
Learning objectives	17
1.1 Why climate science matters for risk assessment	17

1.2	Essential climate science topics	17
	Further reading	18
2	Probability and inference	19
	Learning objectives	19
2.1	Probability theory fundamentals	19
2.1.1	Basic concepts	19
2.1.2	Key distribution functions	20
2.1.3	Multiple variables	21
2.1.4	Examples	22
2.1.5	Summary statistics	23
2.2	Generative models	23
2.2.1	Building models from conditional blocks	24
2.2.2	Visualizing model structure with graphical models	24
2.2.3	Common examples	24
2.3	Inference	24
2.3.1	Transformation of variables	25
2.3.2	The inverse problem and the likelihood function	26
2.3.3	Maximum likelihood estimation	27
2.3.4	Bayesian inference	28
2.4	Computational methods and asymptotic foundations	31
2.4.1	Asymptotic theory	31
2.4.2	Monte Carlo methods	32
2.5	Examples and applications	32
2.5.1	The coin flip (Beta-Binomial model)	32
2.5.2	Linear regression	36
2.5.3	Logistic regression	40
	Further reading	41
3	Machine learning and nonparametric methods	42
	Learning objectives	42
3.1	Essential machine learning concepts	42
	Further reading	42
4	Correlation and dimensionality	43
	Learning objectives	43
4.1	Essential concepts	43
	Further reading	43
5	Model validation and comparison	44
	Learning objectives	44
5.1	Essential model validation concepts	44
	Further reading	44
III	II: Hazard Assessment	45
6	Extreme value theory	46
	See first	46

Learning objectives	46
6.1 Motivation	46
6.2 Approaches for modeling extreme values	48
6.2.1 Block maxima	48
6.2.2 Peaks-over-threshold	48
6.3 Asymptotic theory for extremes	48
6.3.1 Theory for block maxima: The GEV distribution	49
6.3.2 Theory for threshold exceedances: The GPD	49
6.3.3 The shape parameter	50
6.3.4 Connection between GEV and GPD	50
6.4 Return periods and return levels	50
6.4.1 Definitions and Calculation	50
6.4.2 Return level plots	51
6.5 Inference	51
6.5.1 Plotting Positions	51
6.5.2 Moments	52
6.5.3 MLE	52
6.5.4 Bayesian	52
6.6 Sampling variability	52
6.7 Regionalization	52
6.8 Nonstationarity	52
Further reading	52
7 Downscaling and Bias Correction	53
Learning objectives	53
Further reading	53
8 Stochastic weather generators	54
Learning objectives	54
8.1 Essential concepts	54
Further reading	54
9 Physics-based models and calibration	55
Learning objectives	55
9.1 Essential concepts	55
Further reading	55
10 Optimal sampling methods	56
Learning objectives	56
10.1 Essential sampling concepts	56
Further reading	56
11 Global sensitivity analysis	57
Learning objectives	57
11.1 Essential sensitivity analysis concepts	57
Further reading	57

IV III: Risk Management	58
12 Exposure and Vulnerability	59
12.1 Learning objectives	59
12.2 Further reading	59
13 Cost-Benefit Analysis and Net Present Value	60
Learning objectives	60
Further reading	60
14 Policy Search & Optimization	61
15 Risk Transfer	62
See first	62
Learning objectives	62
15.1 Insurance/reinsurance fundamentals	62
15.2 Parametric vs. indemnity coverage	62
15.3 Catastrophe bonds, index-based schemes	62
15.4 Challenges in emerging markets and vulnerable regions	62
Further reading	62
16 Deep Uncertainty and Model Structure	63
See first	63
Learning objectives	63
Further reading	63
17 Robustness	64
18 Adaptive Planning and Flexibility	65
See first	65
Learning objectives	65
Further reading	65
19 Working with People: Values, Participation, and Communication	66
Further reading	66
V Computational Case Studies	67
Overview	68
Probability, inference, and computation: examples	69
Distribution examples	69
Coin flip examples	71
Likelihood visualization	71
Maximum likelihood estimation	72
Bayesian analysis	73
Linear regression example	75
Generating synthetic data	75
Curve fitting by minimizing prediction error	76

Maximum likelihood estimation	77
Bayesian inference	78
MLE with Turing	79
Comparison of approaches	79
Extreme Value Theory Examples	81
Shape parameter implications	81
Houston precipitation data	82
Processing GHCN format	82
Converting dates and units	82
Filling missing dates for complete time series	83
Daily precipitation time series	83
Statistics Without the Agonizing Details	85
Learning Objectives	85
Data	85
A simple analysis	85
The skeptic’s argument	86
Implementation	86
Plotting	86
Alternative	87
References	88
Appendices	90
A Software Setup	90
A.1 Quick start	90
A.1.1 Installation steps	90
A.1.2 Verification	91
A.2 Dig deeper	91
A.2.1 Julia	91
A.2.2 Quarto	91
A.2.3 Visual Studio Code	92
A.2.4 Git and GitHub	92
B Julia Learning Resources	93
B.1 Why Julia?	93
B.2 Learning resources	93
B.2.1 Specialized topics	94
C Large Language Models (“AI”)	95

Welcome

Welcome to Climate Risk Assessment and Management, an online textbook **under construction** by [James Doss-Gollin](#).

! Under Construction

This textbook is a work in progress. Currently it's largely a brain dump, but I am building it out incrementally for use in my own classes. As I add and organize content, I will update the chapter status codes: (*planning*), (*draft*), (*revision*), and (*ready*). [Contributions are welcome!](#)

Motivation and Scope

History

This project emerged from two courses taught at Rice University by [James Doss-Gollin](#): [CEVE 543](#) focused on climate hazard and extremes and [CEVE 421/521](#) focused on risk management.

Aim

The book is motivated by questions like

- What is the probability distribution of wind speeds that a building structure might experience?
- What will the probability distribution of extreme rainfall be in 2050, and what drives uncertainty in this estimate?
- What is the probability distribution of tropical cyclone losses across a regional portfolio?
- When, and how high, should a house be elevated to proactively manage future flood risk?
- What are robust, efficient, and equitable strategies for reducing flood risk in an urban area?

These questions span scales and sectors, yet they share fundamental challenges: characterizing extreme events, quantifying uncertainty, assessing risks, and making robust decisions when probability distributions are unknown or contested. Moreover, there is not a single correct answer to these questions, or a single method that will incontrovertibly answer them.

How to Use This Resource

The book is designed to be useful for practitioners, students, and teachers. Teachers may use individual chapters in their courses. Students may use it as a class text or reference. Practitioners may focus on specific chapters relevant to their work. Each chapter includes learning objectives and can be read independently, though some chapters build on concepts introduced in others.

Structure

- [The Preface](#) introduces the book’s motivation and frames key challenges
- [Part 1](#) introduces key topics in probability, inference, Bayesian methods, optimization, machine learning, and Earth science. Rather than providing a comprehensive treatment, this part focuses on essential concepts and links to further resources.
- [Part 2](#) focuses on hazard **assessment**, namely modeling climate hazards and extremes. Material is organized around thematic applications and predictive tasks. The foundational idea is integrating information from noisy and/or biased sources to estimate the joint probability distribution of relevant hydroclimatic variables.
- [Part 3](#) risk management, which involves both mapping hazard to risk and designing interventions to manage these risks. Key ideas include the sequential nature of decisions, the pursuit of unclear and/or contested objectives, and the need to account for the sensitivity of estimated probability distributions (of hazard and of other relevant physical, social, and economic variables) to underlying models and assumptions.
- [Computational notebooks](#) written in [Julia](#) illustrate and complement the methods and concepts discussed in the text. While notebooks are referenced in the text, they are designed as standalone and self-contained resources.

Prerequisites

Basic probability and multivariate calculus, along with linear algebra, are sufficient mathematical foundations for this textbook. Some exposure to Earth science, hydrology, water resources, or related topics is strongly encouraged for context, though not strictly necessary for understanding methods. This book builds on a wide range of topics and methods in statistics, machine learning, optimization, and Earth science, and expertise in any of these areas may deepen your understanding, but is not necessary. No programming is required to read the book, but going through computational examples and applying methods to your own problems, which can substantially strengthen your understanding, does require programming.

Part I

About this book

License

This textbook is licensed under the [CC BY-NC 4.0 License](#). It is free to use, share, and adapt for non-commercial purposes, provided that you give appropriate credit, provide a link to the license, and indicate if changes were made. If you would like to use this content for commercial purposes, please contact me.

Contributing

This textbook is a work in progress, and we welcome your contributions. Whether it's fixing a typo or proposing a new module, every suggestion helps. The easiest way to contribute is to fork the repository and submit a pull request. If you're not comfortable with that workflow, please open an issue [on GitHub](#).

Citing

Please cite this resource as

```
@book{doss-gollin_textbook:2025,  
  author = {Doss-Gollin, James},  
  title = {Climate Risk Management},  
  year = {2025},  
  url = {https://jdossgollin.github.io/climate-risk-book},  
}
```

In the future, we will move to stable releases with numbered versions.

Further Reading

Climate risk assessment and management are complex and interdisciplinary topics, and we are by no means comprehensive here. This page provides some helpful resources (textbooks, detailed online tutorials, and class websites) for your continued and supplementary study.

Inspiration

This textbook draws inspiration and content from several courses and lecture notes, and I am grateful to the instructors who have shared their materials with me.

- Upmanu Lall’s Environmental Data Analysis course at Columbia
- Vivek Srikrishnan’s [Environmental Systems Analysis](#) and [Climate Risk Analysis](#) classes at Cornell
- R. Balaji’s Advanced Data Analysis Techniques (Statistical Learning Techniques for Engineering and Science) [course](#) at CU Boulder
- Alberto Montanari’s [collection of open course notes and lectures](#)
- **Applegate and Keller (2015)** motivates this project and demonstrates problem-based learning.

Stats + ML basics

This book assumes familiarity with these topics, but these resources may be helpful as a refresher.

- **Blitzstein and Hwang (2019)** provides a thorough introduction to key concepts and ideas in probability. The book accompanies a free online course, [Stat 110](#), which is a great resource for learning probability and statistics. Practice problems and solutions, handouts, and lecture videos are all available online.
- **Downey (2021)** offers an introduction to Bayesian statistics using computational methods. It’s not environment focused but provides code and a clear explanation of core concepts.
- **Gelman (2021)** is a textbook designed for a first course on applied statistics. Clear and well-worked examples underpin discussion of fundamental ideas in statistical analysis and thinking about data.

Applications

There are lots of related books on catastrophe modeling, water resources research, geostats, statistical hydrology and related topics. Here is an incomplete list of some core references.

- **Naghattini (2017)** is a textbook on statistical hydrology that covers many of the same topics as this course. The statistical hydrology literature often obfuscates key ideas with complex notation and terminology, but this book is a helpful introduction to the field.
- **Helsel et al. (2020)** is a comprehensive introduction to water resources and hydrology, focusing on statistical methods for analyzing hydrologic data. Its methods are traditional, with less emphasis on machine learning or Bayesian methods and more attention to null hypothesis significance testing, but its case studies are well-worked and thoughtfully described.
- **Abernathey (2024)** is an excellent resource covering introductory topics in Earth and climate data science using Python, with an emphasis on foundational computations. These core computational concepts serves as a recommended prerequisite for more advanced material in this book.
- **Pyrz (2024)** is a textbook focused on applied machine learning, with a particular focus on geostatistics. There's less focus on extremes, hydroclimate, and decision-making, but it provides very clear and interpretable explanations of many machine learning methods, including some that are not directly covered in this book.
- **Mignan (2024)** is a modern introduction to catastrophe risk modeling that covers a wide range of hazards, including hydroclimatic extremes, from a physics-based perspective. It provides a structured framework for quantifying hazard, exposure, and vulnerability, following industry-standard CAT modeling approaches. While broader in scope and more introductory in level, it complements this book's focus by illustrating foundational principles of probabilistic risk modeling in practice.

More Stats + ML

This book covers a broad set of topics in statistics, machine learning, and optimization. Most chapters could be a textbook of their own, and in fact many exist.

- **Friedman, Hastie, and Tibshirani (2001)** is a classic introduction to machine learning, which complements the Bayesian perspective nicely.
- **Jaynes (2003)** is a classic text on probability theory that you should read if you're interested in questions like "what is probability?"
- **Gelman et al. (2014)** and **McElreath (2020)** are the classic textbooks on Bayesian inference and provide a wealth of insight and detail. The Gelman textbook is a bit more dense while the McElreath book has a more conversational tone, but both cover similar topics.
- **Cressie and Wikle (2011)** provides a detailed exploration of hierarchical space-time models. There have been some computational advances since then that are worth keeping in mind before you apply these models directly, but it's a clearly written and overview.
- **Thurey et al. (2024)** is a new textbook on physics-based deep learning, which is a rapidly growing area of research. It provides a comprehensive overview of the field, including theoretical foundations and practical applications. It covers topics, including neural operators and diffusion models, that are not covered in this course, but which are increasingly used in the climate risk space.
- **Bishop and Bishop (2024)** is a comprehensive, modern, and accessible start-to-finish textbook covering machine learning from basic probability through diffusion models.
- Michael Betancourt's [writing page](#) has detailed and mathematically rigorous explanations of many topics in Bayesian data analysis and probabilistic modeling.

Preface

What is climate risk?

Climate risks arise at the intersection of climate hazards, exposed systems, and vulnerability. They manifest when extreme or changing climate conditions—floods, droughts, extreme temperatures, sea-level rise, or shifting precipitation patterns—impact human and natural systems that are exposed and vulnerable to these conditions. The financial sector terms these “physical risks” to distinguish them from transition risks related to policy and market changes.

Climate risks span scales from the hyperlocal (a single building’s flood exposure) to the global (climate impacts on agricultural productivity). They encompass immediate acute risks from individual extreme events and longer-term chronic risks from gradual climate changes. Crucially, climate risks are not solely natural phenomena but emerge from the complex interactions between climate hazards and the human systems—infrastructure, institutions, communities, and economies—that experience their impacts.

Climate risk is often defined as the product of hazard (probability that something will happen) and consequences (exposure and vulnerability). However, it’s often helpful to start with the decisions we care about.

Risk management

The goal of assessing climate risks is to manage them, as is the focus of Part III. We manage climate risks by

- **building infrastructure**, such as seawalls, stormwater pipes, oyster beds, green roofs, dams
- **designing policy**, such as water pricing, land-use regulations, building codes
- **responding** to climate disasters through disaster response and recovery. While emergency management is beyond the scope of the book, disaster prevention (through infrastructure, policy, etc) and preparation (planning evacuation routes, assessing resource needs, etc) are problems that the tools of this class can inform.

A key insight from considering these applications is that climate risks are not natural phenomena, but occur at the intersection of natural and human systems. A second insight is that decisions about how to manage climate risks do not depend only on climate hazard, but also on human systems and values.

Exposure and vulnerability

Hazards do not create consequences by themselves. Hazards affect things that we care about, whether natural ecosystems, human homes, infrastructure systems, or something else. Quantitatively these are often described as exposure and vulnerability. However, this is not always a helpful framing because everything is exposed, to at least some degree, to climate hazards.

Climate hazard

Climate hazards have several key characteristics:

- **Location-specific impacts:** Specific weather patterns cause different things in different places—tropical cyclones cause extreme winds on the Gulf Coast, while persistent intense rainfall causes flooding in major rivers
- **Require Earth science and data:** Understanding hazards requires both physical process knowledge and empirical data
- **Variable focus on extremes:** Some applications care about extremes, but others (e.g., water management) care about shifts in the whole distribution
- **Multi-scale variability:** Characterized by variability across multiple spatial and temporal scales

What are good strategies?

The simple story

In principle, managing climate risks should be straightforward. If we had clear objectives and well-characterized uncertainty, there are **established mathematical formalisms for decision-making under uncertainty**. Notably, Bayesian Decision Theory provides an elegant framework: find the action a that maximizes expected utility

$$\mathbb{E}[U(a)] = \int U(a, s)p(s)ds,$$

where $U(a, s)$ is the utility of action a given s , and $p(s)$ is the over states of the world. The $\mathbb{E}[U(a)]$ represents the average utility we would expect from action a across all possible future states, weighted by their probabilities (see [Chapter on Probability and Statistics](#) for mathematical foundations).

With this framework and modern advances in operations research and optimization, we could frame climate risk management as a large-scale optimization problem. This might still be a challenging problem, requiring sophisticated optimization methods, large-ensemble Monte Carlo simulation, high-performance computing, and more, but fundamentally **there would be a right answer** that we could identify, at least seek to approximate.

Why this isn't enough

In practice, climate risk management defies this idealized approach for several fundamental reasons:

1. **Deep uncertainty:** Unlike textbook optimization problems, we rarely have well-defined probability distributions over future states. Climate risks involve poorly characterized, multiple, and interacting uncertainties spanning physical processes (climate projections), socioeconomic factors (development patterns, institutional capacity, human behavior), and their complex dependencies. The probability distributions we need span climate hazards, exposure patterns, vulnerability functions, and policy effectiveness—all evolving in ways that resist precise characterization.
2. **Large and poorly defined decision spaces:** The solution space includes not just individual projects but entire systems: infrastructure networks, policy portfolios, risk transfer arrangements, and adaptive management sequences. These decisions interact across scales, sectors, and time horizons in ways that resist comprehensive optimization.
3. **Contested objectives:** Different stakeholders hold different values about what we should optimize for—economic efficiency, equity, robustness, or flexibility. These objectives often conflict, and their relative importance is itself contested and evolving.

This brings us to a crucial insight: **we cannot simply frame climate risk management as a big optimization problem.** The field has witnessed an explosion of computational tools—climate models with ever-finer resolution, machine learning algorithms for processing vast datasets, and sophisticated visualization platforms for rendering complex projections. While these advances represent genuine progress, their proliferation has created new challenges for practitioners seeking to manage real-world climate risks.

The abundance of available tools does not automatically translate to better decisions. Indeed, the sophistication of modern computational approaches can obscure fundamental questions about problem framing, uncertainty characterization, and appropriate methods selection. Without solid conceptual foundations, practitioners may find themselves applying powerful tools inappropriately or mistaking methodological novelty for substantive insight.

The stakes of getting it wrong

The consequences of inadequate climate risk management are severe and diverse. **Infrastructure failures** occur when designs based on historical extremes prove insufficient for future conditions—leading to flooded neighborhoods when storm drains are undersized, or to costly over-design when extreme projections are treated as certainties. **Policy mistakes** compound these problems: development policies that ignore flood risks concentrate vulnerable populations in harm's way, while overly conservative regulations can stifle economic development without commensurate risk reduction benefits.

Financial miscalculations affect both public and private sectors. Insurance companies that underestimate climate risks face catastrophic losses, while those that overestimate risks price themselves out of markets. Infrastructure investors struggle to balance climate resilience against cost constraints, often erring toward solutions that prove either inadequate or prohibitively expensive. These failures cascade across scales: a poorly designed local drainage system contributes to regional

flood management challenges, while flawed national climate risk assessments misguide infrastructure investment priorities across entire countries.

This book

This book develops both the technical tools and conceptual frameworks needed for climate risk management:

- **Part I** provides the statistical, optimization, and machine learning foundations that enable rigorous analysis of climate risks and decision alternatives
- **Part II** focuses on characterizing climate hazards and their uncertainties, emphasizing the integration of multiple imperfect information sources
- **Part III** addresses the transition from hazard to risk and the design of management strategies under deep uncertainty

Throughout, we emphasize that technical sophistication must be coupled with conceptual clarity about the nature of climate risks and the limits of optimization approaches. The goal is not to abandon quantitative analysis, but to use it more wisely—focusing computational power where it adds most value while acknowledging the irreducible uncertainties that require adaptive, robust approaches to climate risk management.

This book aims to teach readers how to **apply** tools from applied mathematics, statistics, and machine learning to answer questions such as

- What is the probability distribution of some relevant hazards or variables, such as (rainfall, wind, flood, temperature, streamflows) at a specific location?
- How do these probability distributions change in the next 50 years?
- How uncertain are these estimates and what specific mechanisms drive these uncertainties?
- What is the distribution of annual losses of a portfolio of assets exposed to one or many climate risks?
- What are trade-offs between up-front costs and future damages for decisions like how high to elevate a house?
- What are robust strategies for sequentially hardening infrastructure against climate risks?
- What are trade-offs between flood and drought protection for managing a reservoir?

While Part I does provide building blocks, they are intended to be self-contained references rather than a comprehensive overview to applied math, statistics, computer science, machine learning, and operations research. Instead, it aims to give you “just enough” context to think carefully about how to apply tools from these fields to climate risk management challenges.

What this book is not

This book focuses on the technical foundations of climate risk assessment and quantitative decision-making under uncertainty. While we address design requirements, social dimensions, and stakeholder considerations throughout—recognizing that technical tools can significantly inform these challenges—there are important aspects of climate risk management that require specialized expertise beyond our scope.

This book will **not** primarily teach you how to:

- **Manage reputational and transition risks:** While we focus on physical climate risks and their quantitative assessment, organizations also face complex risks from changing policies, markets, and stakeholder expectations that require specialized risk management expertise
- **Design and implement adaptive organizations:** While we cover adaptive management strategies and robust decision-making frameworks, the organizational design and management expertise needed to implement these approaches in practice requires additional specialized knowledge
- **Facilitate stakeholder processes:** While the quantitative tools we teach can strongly support consensus building by clarifying trade-offs and uncertainties, the facilitation, negotiation, and collaborative governance skills needed to lead stakeholder processes require specialized training
- **Develop communication strategies:** While we emphasize how to interpret and present quantitative risk assessments, developing effective communication strategies for diverse audiences—policymakers, communities, investors—requires specialized expertise in science communication and public engagement
- **Navigate implementation challenges:** While we address policy design and infrastructure planning from an analytical perspective, the practical challenges of construction management, regulatory processes, and community engagement require domain-specific expertise

This is an interdisciplinary text that draws insights from multiple fields and acknowledges the social, political, and institutional contexts that shape climate risk management. However, our primary focus remains on the quantitative and analytical foundations that can inform—but not replace—the broader expertise needed for effective practice.

Part II

I: Foundations

1 Climate science

Learning objectives

After reading this chapter, you should be able to:

- Understand why climate science is essential for risk assessment
- Identify key climate science topics you should study for effective risk management
- Navigate to appropriate resources for learning climate science fundamentals

1.1 Why climate science matters for risk assessment

Climate hazards don't occur in isolation. A hurricane's intensity depends on sea surface temperatures and atmospheric conditions. Droughts emerge from large-scale circulation patterns and ocean-atmosphere interactions. Floods reflect not just local rainfall but also broader weather systems and seasonal cycles.

Understanding these connections is crucial for risk assessment because climate science helps us answer three fundamental questions:

1. **What physical processes create hazardous weather patterns?** Understanding the mechanisms behind hurricanes, droughts, heat waves, and floods helps us identify when and where they're most likely to occur.
2. **How do these patterns vary naturally over time?** Climate systems exhibit variability on multiple timescales—from seasonal cycles to multi-decadal oscillations—that affect the frequency and intensity of extreme events.
3. **How might climate change alter these patterns?** As greenhouse gas concentrations rise, the statistical properties of weather and climate are shifting, requiring us to account for non-stationarity in our risk assessments.

1.2 Essential climate science topics

1. Climate models and modeling
2. Multi-scale variability
3. Specific weather patterns
4. Climate change and sensitivity
5. Hydrologic cycle

Further reading

- Mudelsee (2020): statistical approaches to climate extremes
- Merz et al. (2014): flood risk methods connecting climate to impacts
- Ghil et al. (2011): physical processes and extreme value behavior

2 Probability and inference

This chapter covers the fundamental concepts of probabilistic modeling and statistical inference. Data analysis proceeds by building a generative model—a formal, probabilistic hypothesis about how data are created. Inference is the inverse problem of using observed data to learn about model parameters. Computational methods make complex inference problems tractable.

Learning objectives

After reading this chapter, you should be able to:

- Build generative models using probability distributions.
- Apply maximum likelihood and Bayesian inference to estimate parameters.
- Use Monte Carlo methods when analytical solutions don't exist.
- Recognize when computational methods are required versus analytical approaches.

2.1 Probability theory fundamentals

The concepts in this section are essential building blocks for generative models. These mathematical tools support practical modeling applications.

2.1.1 Basic concepts

2.1.1.1 Random variables

A random variable is a function that assigns numerical values to the outcomes of a random experiment. Random variables provide the mathematical foundation for describing uncertainty.

2.1.1.2 Notation conventions

- Random variables: Capital letters (X, Y, Z)
- Realizations (specific values): Lowercase letters (x, y, z)
- Parameters: Greek letters (θ, μ, σ)
- Observed data: y (following Bayesian convention)
- Predictions: \tilde{y} (y-tilde)

2.1.1.3 Types of random variables

- Discrete random variables take on countable values (e.g., number of floods per year)
- Continuous random variables take on uncountable values (e.g., temperature, precipitation amount)

The type of random variable determines which mathematical tools we use to describe its behavior.

2.1.2 Key distribution functions

2.1.2.1 Key functions

The foundation of probability theory rests on three fundamental functions that describe random variables.

2.1.2.2 Probability Mass Function (PMF)

For discrete random variables, $P(X = x)$ gives the probability that the variable takes on a specific value x . The PMF satisfies $\sum_x P(X = x) = 1$.

2.1.2.3 Probability Density Function (PDF)

For continuous random variables, $p(x)$ describes the relative likelihood of different values. Crucially, $p(x)$ is not a probability but a density. Since probability is density multiplied by a (potentially very small) interval of x , the value of $p(x)$ itself can exceed 1 without violating the laws of probability. Probabilities are areas under the curve: $P(a \leq X \leq b) = \int_a^b p(x) dx$.

2.1.2.4 Cumulative Distribution Function (CDF)

$F(x) = P(X \leq x)$ gives the probability that a random variable is less than or equal to x . The CDF is the most fundamental descriptor, defined for all random variables (discrete and continuous). It unifies probability concepts and is essential for quantiles and return periods.

2.1.2.5 Quantile function

The quantile function $Q(p) = F^{-1}(p)$ is the inverse of the CDF. It takes a probability $p \in [0, 1]$ and returns the value x such that $P(X \leq x) = p$.

- The median is $Q(0.5)$
- The N -year return level in extreme value analysis is $Q(1 - 1/N)$
- Quantiles enable inverse transform sampling for Monte Carlo simulation

The quantile function has direct applications in climate risk assessment. For example, the 100-year flood has return period 100, meaning it has probability 0.01 in any given year, corresponding to quantile $Q(0.99)$.

2.1.3 Multiple variables

2.1.3.1 Joint, marginal, and conditional distributions

Real systems involve multiple random variables, requiring tools to describe their relationships. This machinery allows construction of complex models from simpler components.

2.1.3.2 Joint Distribution

$p(x, y)$ for continuous or $P(X = x, Y = y)$ for discrete gives the probability of events occurring together.

2.1.3.3 Marginal Distribution

$p(x)$ or $P(X = x)$ gives the probability of an event, irrespective of other variables. Calculated by summing or integrating over the other variables: $p(x) = \int p(x, y) dy$.

2.1.3.4 Conditional Distribution

$p(y | x)$ or $P(Y = y | X = x)$ gives the probability of an event given that another event has occurred. Conditional distributions describe how variables depend on each other.

2.1.3.5 Independence

Two random variables X and Y are independent if their joint distribution is the product of their marginal distributions: $p(x, y) = p(x)p(y)$ for continuous variables, and $P(X = x, Y = y) = P(X = x)P(Y = y)$ for discrete variables. For example, temperature and rainfall on a given day are typically **not** independent—hot days often have lower rainfall probability.

2.1.3.6 IID (Independent and Identically Distributed)

A sequence of random variables that are independent and have the same distribution.

2.1.3.7 Bayes' rule

The mechanical relationship between joint, marginal, and conditional distributions:

$$p(y | x) = \frac{p(x | y)p(y)}{p(x)}$$

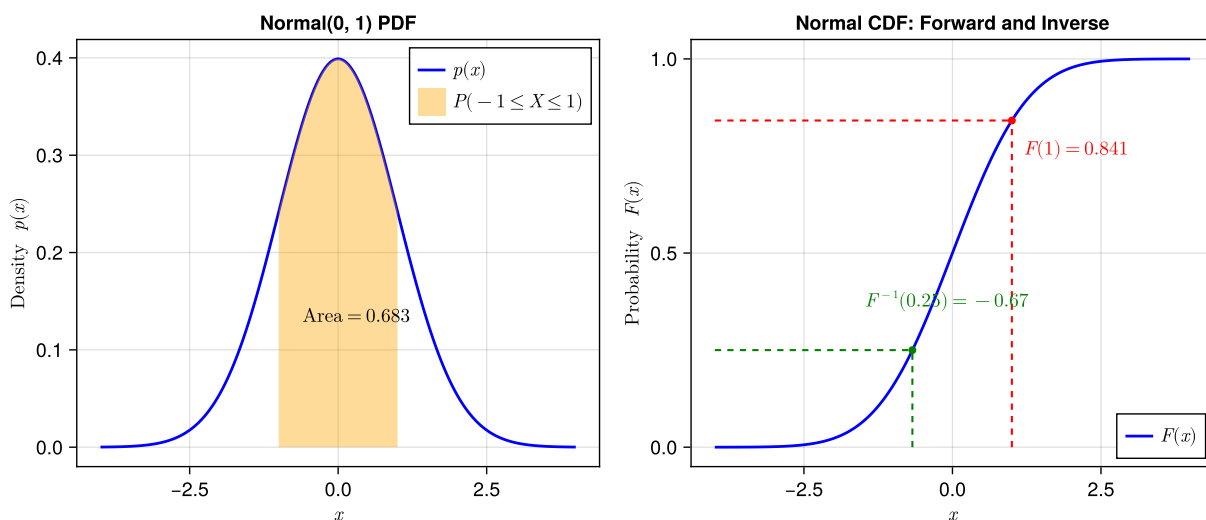
Bayes' rule is a consequence of the definition of conditional probability. It becomes a tool for inference when interpreted probabilistically.

2.1.4 Examples

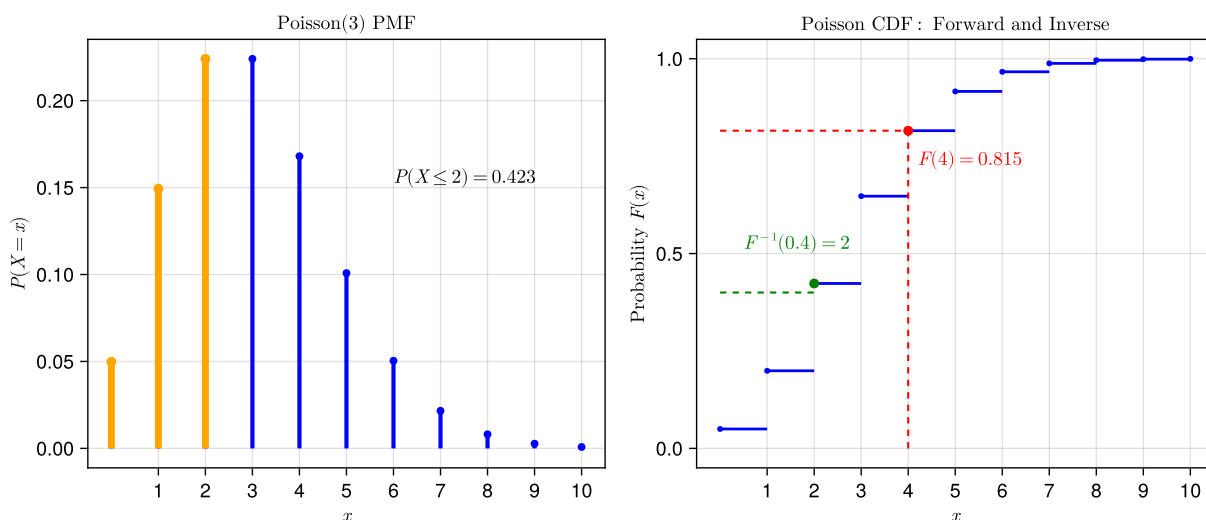
2.1.4.1 Distribution examples

These examples illustrate the key concepts of probability density/mass functions, cumulative distribution functions, and quantiles using two fundamental distributions. Each example shows both the forward operation (finding probabilities from values) and the inverse operation (finding values from probabilities).

The normal distribution demonstrates these concepts for continuous random variables, where probabilities correspond to areas under smooth curves. The Poisson distribution shows the analogous concepts for discrete random variables, where probabilities correspond to point masses and CDFs are step functions.



Next, a Poisson distribution example for discrete random variables.



2.1.5 Summary statistics

2.1.5.1 Key descriptors: expectation and moments

Probability distributions are completely described by their PDF/PMF and CDF, but we often need summary statistics that capture essential properties.

2.1.5.2 Expectation (Expected Value)

The expectation is the formal definition of the quantity we approximate with the sample mean in a Monte Carlo simulation. The expectation of a function $g(X)$ is:

$$\mathbb{E}[g(X)] = \int g(x)p(x) dx$$

for continuous variables, or $\mathbb{E}[g(X)] = \sum_x g(x)P(X = x)$ for discrete variables.

2.1.5.3 Moments of a distribution

- Mean: $\mu = \mathbb{E}[X]$ measures central tendency
- Variance: $\sigma^2 = \mathbb{E}[(X - \mu)^2]$ measures spread or scale; standard deviation is $\sigma = \sqrt{\sigma^2}$
- Higher-order moments: Skewness measures asymmetry, kurtosis measures tail weight

For heavy-tailed distributions used in extreme value theory, some higher-order moments (or even the variance) may not exist because the defining integrals diverge.

2.2 Generative models

The probability theory we've covered provides the mathematical language for uncertainty. Now we turn to its primary application: building models that describe how data are created. These generative models form the foundation for statistical inference—the process of learning from data.

Statistical modeling proceeds by creating a forward model—a formal, probabilistic description of how data are created. These models serve two purposes for scientific inference and decision-making.

First, they function as simulation engines. A generative model enables simulation of data under different parameter settings. This supports uncertainty propagation through complex systems, model checking by comparing simulated and observed data, and exploration of model assumptions.

Second, they provide a systematic framework for building complex models from simple pieces. A complicated joint distribution can be decomposed into a sequence of simpler conditional distributions. This modular approach supports hierarchical modeling and incorporation of domain knowledge at each stage.

2.2.1 Building models from conditional blocks

A complex joint distribution can be decomposed into a sequence of simpler conditional distributions. This decomposition constitutes the model structure.

The chain rule of probability (this is where Bayes' theorem comes from) tells us that any joint distribution can be factored as:

$$p(A, B) = p(A) \cdot p(B \mid A)$$

or, for three variables:

$$p(A, B, C) = p(A) \cdot p(B \mid A) \cdot p(C \mid A, B)$$

2.2.2 Visualizing model structure with graphical models

We can represent the conditional dependence structure of our generative story visually using a Directed Acyclic Graph (DAG).

Nodes represent random variables (parameters or data). Arrows represent conditional dependence. An arrow from A to B means that the value of B is generated based on the value of A.

This provides a clear, unambiguous map of the model's assumptions. DAGs support model design, communication, and debugging.

2.2.3 Common examples

Examples of generative models include:

- **Signal + Noise:** The classic model, e.g., $y = f(x; \beta) + \epsilon$, where we specify a distribution for the noise term ϵ .
- **Dynamical Systems:** A system of Ordinary Differential Equations (ODEs) with unknown parameters that govern the system's evolution.
- **Difference Equations:** Models of processes over discrete time, potentially including random walks.

2.3 Inference

Forward modeling simulates data given parameters. Inference solves the inverse problem: finding parameters that explain observed data. The forward direction is direct: given parameters, simulate data. The inverse direction requires mathematical machinery to work backwards from data to parameters.

2.3.1 Transformation of variables

A core task in probabilistic modeling is understanding how randomness propagates through a system. If we have a random variable X with a known probability density function (PDF), $p_X(x)$, and we create a new random variable Y by applying a function, $Y = g(X)$, what is the PDF of Y , denoted $p_Y(y)$?

It's tempting to just substitute x with $g^{-1}(y)$ in the original PDF. This is incorrect because the function g can stretch or compress the space of outcomes. We must account for this distortion to ensure the total probability remains 1.

2.3.1.1 Using the cumulative distribution function

The CDF of the new variable relates to the original CDF through a three-step process.

The most reliable way to solve this is to start with the CDF, $F(x) = P(X \leq x)$, because probabilities are always conserved. Assume that $Y = g(X)$ is a strictly increasing function (like $Y = e^X$).

By definition, the CDF of Y is $F_Y(y) = P(Y \leq y)$.

1. Substitute the function: $F_Y(y) = P(g(X) \leq y)$
2. Apply the inverse function: Since g is increasing, we can apply its inverse, g^{-1} , to both sides of the inequality without flipping the sign: $F_Y(y) = P(X \leq g^{-1}(y))$
3. Recognize the original CDF: The expression on the right is the CDF of X evaluated at the point $g^{-1}(y)$: $F_Y(y) = F_X(g^{-1}(y))$

2.3.1.2 From CDF to PDF: the final step with the chain rule

To find the PDF, we differentiate the CDF with respect to y . Applying the chain rule to the expression above gives us our answer:

$$p_Y(y) = \frac{d}{dy} F_Y(y) = \frac{d}{dy} F_X(g^{-1}(y))$$

$$p_Y(y) = p_X(g^{-1}(y)) \cdot \frac{d}{dy} g^{-1}(y)$$

If $g(X)$ were a decreasing function, we would get a negative sign, which is handled by taking the absolute value. This gives us the general change of variables formula:

$$p_Y(y) = p_X(g^{-1}(y)) \left| \frac{d}{dy} g^{-1}(y) \right|$$

2.3.1.3 Conservation of probability mass

Think of the PDF as describing how a unit of “probability mass” is spread along the x -axis. The total area under the curve must be 1. When we apply the function $Y = g(X)$, we are stretching and compressing the axis itself.

- If the function stretches a region, the probability density must decrease proportionally to keep the area constant.
- If the function compresses a region, the density must increase.

The term $\left| \frac{d}{dy} g^{-1}(y) \right|$ is the Jacobian of the transformation. It is precisely the “stretching factor” that tells us how much the coordinate system is distorted at each point, ensuring our total probability mass is conserved.

For more on this somewhat unintuitive point, see section 1.8 of Gelman et al. (2014) or [this blog post](#).

2.3.2 The inverse problem and the likelihood function

The central tool for connecting data to parameters is the likelihood function. The likelihood is the conditional probability $p(y \mid \theta)$, where y represents our observed data.

2.3.2.1 Definition

The likelihood tells us how likely we are to see the observed data y for some value of the parameters θ .

2.3.2.2 Crucial distinction

The likelihood is not the probability of the parameters. It’s the probability (or probability density) of the data given the parameters.

This confusion is common: $p(\text{data}|\text{parameters})$ tells us about data likelihood, not parameter probability. Only Bayesian inference gives us $p(\text{parameters}|\text{data})$.

For continuous variables, since we’re dealing with a density, the probability of getting exactly that value is zero, but the probability of getting near it is the integral of the PDF over a small interval.

2.3.2.3 Independence and the product form

If we assume our data points are independent and identically distributed (IID), then by the definition of independence:

$$p(y_1, y_2, \dots, y_n \mid \theta) = \prod_{i=1}^n p(y_i \mid \theta)$$

2.3.2.4 The log-likelihood

Products are numerically unstable and difficult to work with. Since the logarithm is monotonic, $\max_{\theta} p(y | \theta)$ has the same solution as $\max_{\theta} \log p(y | \theta)$. For IID data, this gives us:

$$\log p(y | \theta) = \sum_{i=1}^n \log p(y_i | \theta)$$

2.3.3 Maximum likelihood estimation

Maximum likelihood estimation (MLE) finds the parameter values that maximize the likelihood function:

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} p(y | \theta)$$

2.3.3.1 Why maximum likelihood makes sense

The likelihood function $p(y | \theta)$ gives the probability of observed data under different parameter values. Maximum likelihood estimation finds the parameter values that maximize the probability of the observed data. The approach selects parameters that best explain the observations.

In climate applications, MLE estimates parameters of distributions describing rainfall, temperature, or wind speed by finding values that maximize the probability of historical observations. The estimates inform risk assessments and infrastructure design.

2.3.3.2 Implementation

We find the actual parameter values using optimization approaches. This may involve analytical differentiation (setting derivatives to zero) or numerical optimization methods when closed-form solutions don't exist.

This reframes the statistical problem of inference as a numerical problem of optimization.

2.3.3.3 Properties and theoretical foundations

Understanding when and why MLE works requires defining estimator quality. An estimator should be consistent (converge to the true value as sample size increases), efficient (achieve low variance), and unbiased (correct on average).

Under regularity conditions, MLE estimators have desirable asymptotic properties. As the sample size n grows large, the MLE estimator $\hat{\theta}_{\text{MLE}}$ becomes consistent—it converges to the true parameter value θ_0 . The estimator also becomes asymptotically normal, meaning its sampling distribution approaches a normal distribution centered on the true value. Finally, MLE achieves the Cramér-Rao lower bound, making it asymptotically efficient among unbiased estimators.

The Fisher Information $I(\theta) = -\mathbb{E} \left[\frac{d^2}{d\theta^2} \log p(Y | \theta) \right]$ quantifies how much information the data contains about the parameter. Higher Fisher Information means the likelihood function is more peaked, leading to more precise parameter estimates.

2.3.3.4 The assumption of “true” parameters

These theoretical properties assume there exists a single “true” parameter value θ_0 that generated our data, and that our model is correctly specified. But in climate science, we often use simplified models to describe extremely complex phenomena. Is there really a single “true” rainfall parameter for a location when climate is nonstationary, multiple physical processes interact, and our models are necessarily approximations of reality?

This question motivates the Bayesian approach to inference. Bayesian methods acknowledge that parameter values are uncertain and provide probability distributions that quantify this uncertainty. The Bayesian perspective is valuable for complex, imperfect models of climate systems.

2.3.3.5 Computational considerations

Finding maximum likelihood estimates requires different approaches depending on the complexity of the model.

Analytical solutions exist when we can solve $\frac{d}{d\theta} \log p(y | \theta) = 0$ in closed form. This works for simple models like Normal distributions with known variance, or the coin flip example we examine below. These cases provide valuable intuition and serve as building blocks for more complex problems. Additional computational implementations of these methods are provided in the [companion notebook](#).

Numerical optimization becomes necessary when no closed-form solution exists. Common algorithms include gradient-based methods like Newton-Raphson and BFGS, which use derivative information to efficiently locate the maximum. The Expectation-Maximization (EM) algorithm handles models with latent variables by iteratively estimating missing data and updating parameters. Grid search remains useful for low-dimensional problems where visualization of the likelihood surface provides insight.

Practical challenges arise in applications. The likelihood surface may contain multiple local maxima, requiring different starting values to find the global optimum. Numerical stability requires working with log-likelihoods rather than products of small probabilities. Flat likelihood surfaces indicate that data contain limited information about parameters. All methods assume correct model specification—poor model approximations yield misleading results regardless of optimization quality.

2.3.4 Bayesian inference

Bayesian inference takes a fundamentally different approach to the inverse problem. Rather than finding a single “best” parameter value, we use the data to update our beliefs about the parameters, resulting in a full probability distribution that naturally quantifies uncertainty.

The Bayesian interpretation of Bayes’ rule:

$$p(\theta | y) = \frac{p(y | \theta)p(\theta)}{p(y)}$$

where $p(\theta)$ is our **prior** belief about the parameters before seeing the data, $p(y | \theta)$ is the **likelihood** (the same function used in maximum likelihood estimation), $p(\theta | y)$ is the **posterior** distribution representing our updated beliefs, and $p(y)$ is a normalizing constant ensuring the posterior integrates to 1.

Since $p(y)$ doesn't depend on θ for a fixed dataset, we often write:

$$p(\theta | y) \propto p(y | \theta)p(\theta)$$

The posterior distribution is the primary output of Bayesian analysis. It tells us not just what parameter values are most plausible, but also how uncertain we should be about them. However, computing this posterior distribution requires different approaches depending on the complexity of the model.

2.3.4.1 Maximum A Posteriori (MAP) estimation

The simplest approach to Bayesian inference is to find the mode of the posterior distribution—the parameter value that is most probable given the data. This is called Maximum A Posteriori (MAP) estimation:

$$\hat{\theta}_{MAP} = \arg \max_{\theta} p(\theta | y) = \arg \max_{\theta} p(y | \theta)p(\theta)$$

Since the logarithm is monotonic, this is equivalent to:

$$\hat{\theta}_{MAP} = \arg \max_{\theta} [\log p(y | \theta) + \log p(\theta)]$$

MAP estimation provides a natural bridge between maximum likelihood and full Bayesian inference. When the prior is uniform (non-informative), MAP reduces to maximum likelihood estimation. When the prior is informative, it regularizes the estimate, preventing overfitting.

However, MAP gives only a point estimate and doesn't quantify uncertainty. To get the full benefit of the Bayesian approach, we need the entire posterior distribution.

2.3.4.2 Analytic posterior (conjugate priors)

In special cases, we can compute the posterior distribution analytically using conjugate priors. A prior is conjugate to a likelihood if the posterior has the same functional form as the prior. This mathematical convenience allows us to update our beliefs with a simple algebraic formula.

Conjugacy provides exact formulas for belief updating. The posterior parameters are the prior parameters plus contributions from the data. As data accumulate, the prior influence diminishes relative to the likelihood.

Conjugate priors exist for several important distributions, including exponential family distributions like the Normal, Binomial, and Poisson. However, conjugate priors exist for only a limited set of models. Most real-world problems require computational approaches.

2.3.4.3 Computational posterior (MCMC)

For most real problems, the posterior distribution cannot be computed analytically. The denominator $p(y)$ requires integrating over all possible parameter values, which is often intractable. Markov Chain Monte Carlo (MCMC) methods become essential.

MCMC algorithms create a Markov chain whose stationary distribution is the target posterior distribution. Instead of computing the posterior directly, we simulate from it by drawing samples $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)}$. These samples allow us to approximate any posterior quantity of interest:

$$\mathbb{E}[g(\theta) \mid y] \approx \frac{1}{N} \sum_{i=1}^N g(\theta^{(i)})$$

Common MCMC algorithms include:

- **Metropolis-Hastings:** A general-purpose algorithm that uses a proposal distribution and an acceptance criterion
- **Gibbs sampling:** Efficient for multivariate problems where conditional distributions are available
- **Hamiltonian Monte Carlo:** Uses gradient information for more efficient exploration of the parameter space

MCMC requires only the posterior up to a constant of proportionality: $p(\theta \mid y) \propto p(y \mid \theta)p(\theta)$. This enables MCMC for complex models where computing the normalizing constant is impossible.

MCMC diagnostics assess whether chains have converged to the target distribution. Poor convergence leads to biased estimates, making diagnostic checking essential for reliable inference.

2.3.4.4 Uncertainty estimates: definitions

2.3.4.4.1 Credible interval (Bayesian)

For a parameter θ with posterior density $p(\theta \mid y)$, a $(1 - \alpha)$ credible set $C(y)$ satisfies

$$\int_{C(y)} p(\theta \mid y) d\theta = 1 - \alpha.$$

For a univariate θ , a common choice is the equal-tailed credible interval

$$[Q_{\theta|y}(\alpha/2), Q_{\theta|y}(1 - \alpha/2)],$$

where Q denotes the posterior quantile function. Another choice is the highest posterior density (HPD) interval, the shortest interval whose posterior probability is $1 - \alpha$. Credible intervals are probability statements about θ conditional on the observed data (and prior) (e.g., Gelman et al. 2014; Downey 2021).

2.3.4.4.2 Confidence interval (frequentist)

A $(1 - \alpha)$ confidence set is a data-dependent rule $C(Y)$ such that, for all fixed θ ,

$$\Pr_{Y \sim p(\cdot|\theta)} \{ \theta \in C(Y) \} \geq 1 - \alpha.$$

In words, if we were to repeat the experiment many times under the same θ , the procedure would cover the true θ at least $1 - \alpha$ of the time. The probability is over hypothetical repetitions of the data, not over θ . In simple models, confidence intervals arise from pivots or large-sample normal approximations (see, e.g., Gelman 2021).

2.3.4.5 Credible intervals vs. confidence intervals

Practitioners often report “intervals,” but Bayesian credible intervals and frequentist confidence intervals answer different questions.

- Credible interval (Bayesian): An interval C such that $P(C \mid y) = 1 - \alpha$. It is a probability statement about the parameter given the observed data and prior. Common choices are equal-tailed intervals and highest posterior density (HPD) intervals. Results depend on both the likelihood and the prior.
- Confidence interval (frequentist): A procedure that, under repeated sampling from the data-generating process, covers the true (fixed) parameter at rate $1 - \alpha$. After observing one dataset, the interval either contains the true value or it does not; the $1 - \alpha$ statement refers to the long-run frequency of the procedure, not the probability for this particular dataset.

The two intervals can coincide in simple regular problems (e.g., large samples with weak/flat priors), but they generally differ in finite samples or with informative priors. Be careful not to interpret a 95% confidence interval as “95% probability that θ lies in the interval” for a specific dataset—that interpretation is Bayesian and corresponds to a 95% credible interval.

For clear discussions, see Gelman et al. (2014); Gelman (2021); Downey (2021).

2.4 Computational methods and asymptotic foundations

2.4.1 Asymptotic theory

Two fundamental theorems justify statistical estimation and Monte Carlo methods:

2.4.1.1 Law of Large Numbers

As sample size N grows, the sample mean converges to the true expected value. This theorem underlies both parameter estimation and Monte Carlo methods.

2.4.1.2 Central Limit Theorem

The distribution of a sample mean approaches a Normal distribution as the sample size increases. This enables uncertainty quantification and justifies confidence intervals and bootstrap methods.

2.4.2 Monte Carlo methods

Most decision-relevant quantities can be expressed as expectations. Simulation approximates expectations when analytical calculation is difficult.

2.4.2.1 Simulation and expectations

The basic Monte Carlo estimate of an expectation is:

$$\mathbb{E}[X] \approx \frac{1}{N} \sum_{i=1}^N x_i$$

where x_1, x_2, \dots, x_N are independent samples from the distribution of X .

The Monte Carlo standard error is approximately σ/\sqrt{N} , where σ is the standard deviation of the quantity being estimated.

2.4.2.2 Markov Chain Monte Carlo (MCMC) for Bayesian inference

For most real problems, the posterior distribution cannot be computed analytically. MCMC methods create a Markov chain whose stationary distribution is the posterior distribution. Common algorithms include Metropolis-Hastings and Gibbs sampling. MCMC diagnostics help us assess whether our chains have converged to the target distribution.

2.4.2.3 Bootstrap methods for frequentist inference

Bootstrap methods quantify uncertainty around maximum likelihood estimates. Bootstrap resampling provides a Monte Carlo approach to estimate sampling distributions of estimators.

2.5 Examples and applications

These examples demonstrate progression from analytical solutions to computational methods for complex problems. Computational implementations are provided in the [companion notebook](#).

2.5.1 The coin flip (Beta-Binomial model)

A series of coin flips are independent Bernoulli trials with fixed probability of heads, θ . Given y heads in n flips, the goal is to infer θ .

2.5.1.1 Likelihood

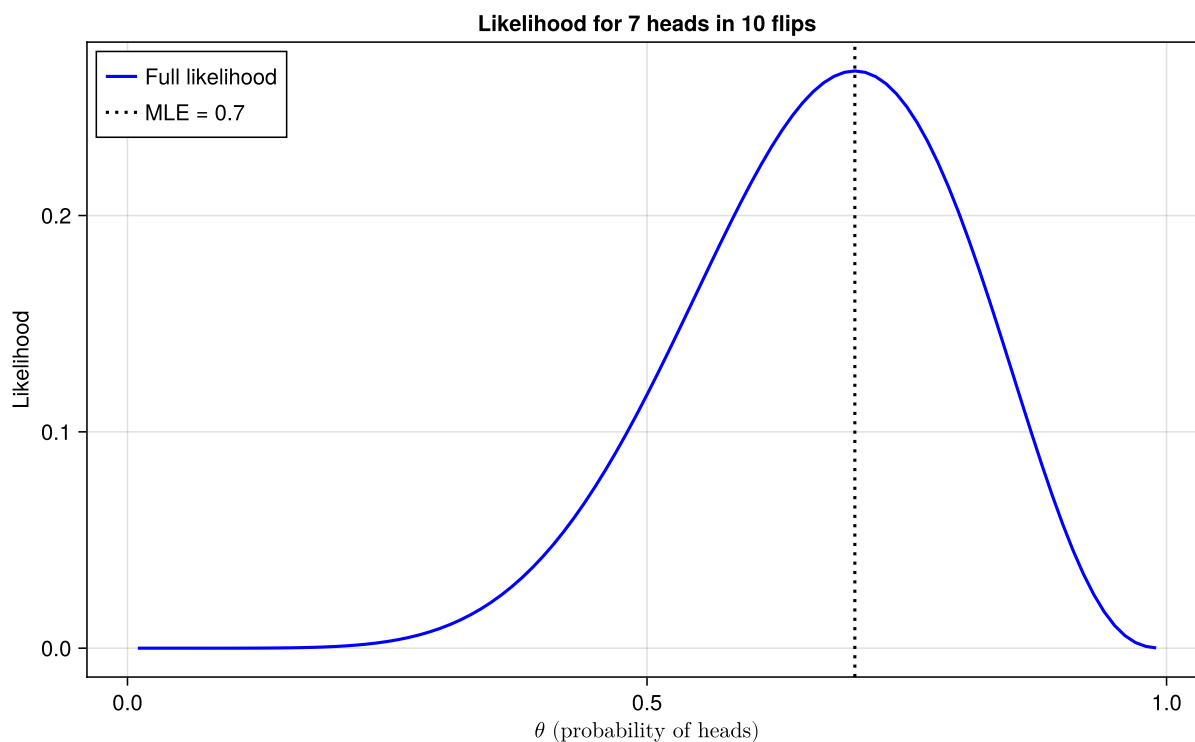
For n independent coin flips with probability θ of heads, if we observe y heads, the likelihood is given by the Binomial distribution:

$$p(y \mid \theta, n) = \binom{n}{y} \theta^y (1 - \theta)^{n-y}$$

where $\binom{n}{y}$ is the binomial coefficient. Since we condition on the observed data y and n , the binomial coefficient is a constant, so for inference purposes we can work with the proportional likelihood:

$$p(y \mid \theta) \propto \theta^y (1 - \theta)^{n-y}$$

The following implementation visualizes how the likelihood varies with θ :



The likelihood peaks at the observed proportion of heads, $\hat{\theta}_{MLE} = y/n$, though many other values of θ remain plausible given limited data. A value of $\theta = 0.1$ would make the observed data very unlikely.

2.5.1.2 Maximum likelihood solution

The maximum likelihood estimate can be derived analytically by maximizing the log-likelihood:

$$\log p(y \mid \theta) = \log \left(\binom{n}{y} \theta^y (1 - \theta)^{n-y} \right)$$

which simplifies to:

$$\log p(y | \theta) = y \log \theta + (n - y) \log(1 - \theta) + \text{const.}$$

We don't need to keep track of the constant because it doesn't depend on θ . Taking the derivative with respect to θ and setting to zero:

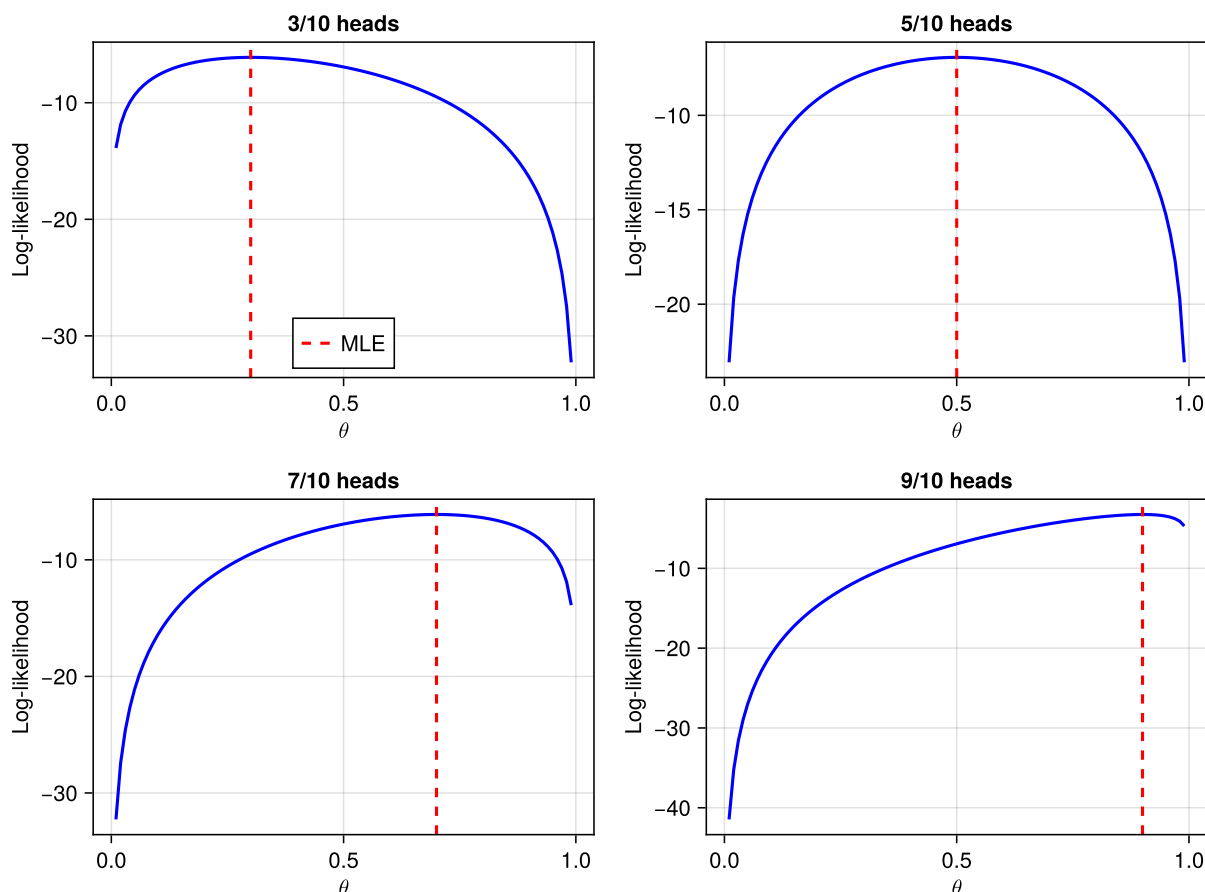
$$\frac{d}{d\theta} \log p(y | \theta) = \frac{y}{\theta} - \frac{n - y}{1 - \theta} = 0$$

Solving for θ :

$$\frac{y}{\theta} = \frac{n - y}{1 - \theta} \implies y(1 - \theta) = (n - y)\theta \implies y = n\theta$$

Therefore: $\hat{\theta}_{MLE} = \frac{y}{n}$

The maximum likelihood estimate is the observed proportion of heads. This example demonstrates the general MLE procedure: write down the likelihood function, take the log, differentiate with respect to the parameters, set the derivative to zero, and solve for the parameter values. This analytical approach works for simple models and provides the template for more complex problems that require numerical optimization. This pattern holds across different experimental outcomes:



2.5.1.3 Bayesian solution

The Beta distribution is conjugate to the Binomial likelihood. That means that if we choose a Beta prior for θ , the posterior will also be a Beta distribution.

Prior: Our prior is $\theta \sim \text{Beta}(\alpha, \beta)$ with density:

$$p(\alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} \theta^{\alpha-1} (1 - \theta)^{\beta-1}$$

The likelihood is as before:

$$p(y | \theta) = \binom{n}{y} \theta^y (1 - \theta)^{n-y}$$

Using Bayes' rule:

$$p(\theta | y) \propto p(y | \theta) p(\theta)$$

Substituting and collecting terms:

$$p(\theta | y) \propto \theta^y (1 - \theta)^{n-y} \cdot \theta^{\alpha-1} (1 - \theta)^{\beta-1}$$

so

$$p(\theta | y) \propto \theta^{(y+\alpha)-1} (1 - \theta)^{(n-y+\beta)-1}$$

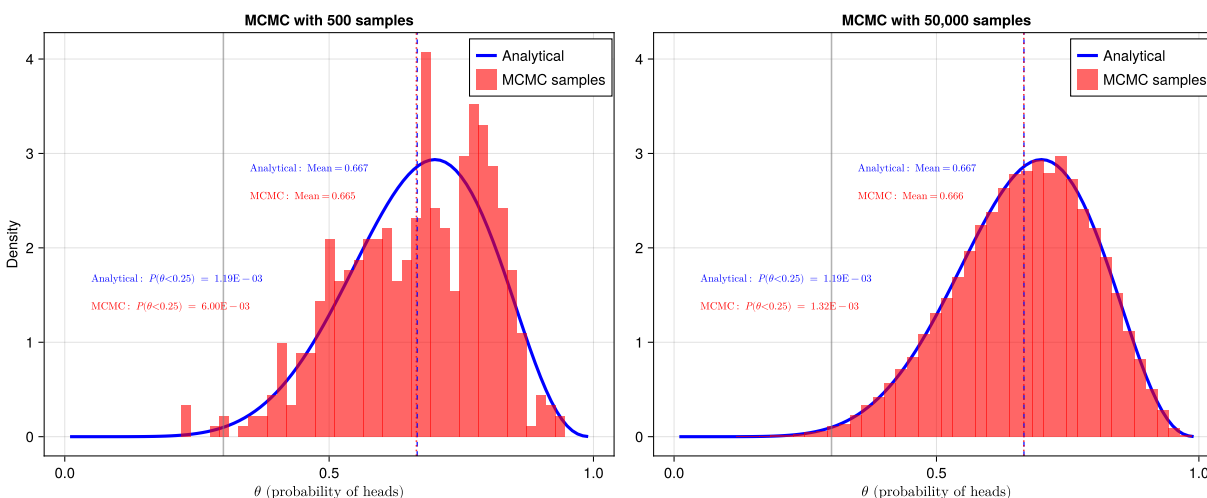
This has the form of a Beta distribution! Therefore:

$$\theta | y \sim \text{Beta}(\alpha + y, \beta + n - y)$$

The posterior parameters are the prior parameters plus contributions from the data:

- α (prior “successes”) + y (observed successes)
- β (prior “failures”) + $(n - y)$ (observed failures)

The following compares the analytical solution with MCMC sampling using Turing.jl:



The analytical and MCMC results match closely, demonstrating that the Beta-Binomial conjugacy gives exact results, MCMC provides a general computational approach when analytical solutions don't exist, and both approaches quantify uncertainty through the full posterior distribution.

MCMC accuracy varies across the distribution. The method provides better estimates in the center of the distribution than in the tails. For example, computing $\mathbb{E}[\theta]$ requires fewer samples than computing the probability that $\theta \leq 0.3$ (formally, $\mathbb{E}[\mathbf{1}(\theta < 0.25)]$ where $\mathbf{1}$ is the indicator function), since the latter depends on accurately characterizing the probability mass in the tail. This reflects a general principle: quantities that depend on rare events or extreme values require more samples and/or more sophisticated sampling strategies for accurate estimation.

Just as we infer from coin flips, we estimate the probability of extreme rainfall from historical data. The same statistical principles apply whether we're analyzing controlled experiments or climate observations.

2.5.2 Linear regression

Given n observations of response variable y_i and predictor variable x_i , the goal is to understand the relationship between x and y . Linear regression principles apply directly to modeling relationships like temperature trends over time or the relationship between atmospheric CO₂ and global temperature.

2.5.2.1 Generative story

The response follows a linear relationship with added noise:


$$y_i = a + bx_i + \epsilon_i$$

where:

- a is the intercept (value of y when $x = 0$)
- b is the slope (change in y per unit change in x)
- $\epsilon_i \sim \text{Normal}(0, \sigma^2)$ is independent random noise

This says that each observation is the “true” linear relationship plus some random variation.

The following shows the synthetic dataset used throughout the examples:



2.5.2.2 Curve fitting by minimizing loss

Without a probabilistic model, this becomes a curve fitting problem: find the line that “best fits” the data.

This requires a loss function to measure estimation quality. A natural choice is the mean squared error (MSE):

$$\text{MSE}(a, b) = \frac{1}{n} \sum_{i=1}^n (y_i - (a + bx_i))^2$$

MSE measures the average squared distance between predictions $(a + bx_i)$ and observed values y_i .

Minimizing this loss function yields the best-fitting line:

$$(\hat{a}, \hat{b}) = \arg \min_{a, b} \sum_{i=1}^n (y_i - (a + bx_i))^2$$

Taking partial derivatives and setting to zero:

$$\frac{\partial}{\partial a} \sum_{i=1}^n (y_i - a - bx_i)^2 = -2 \sum_{i=1}^n (y_i - a - bx_i) = 0$$

$$\frac{\partial}{\partial b} \sum_{i=1}^n (y_i - a - bx_i)^2 = -2 \sum_{i=1}^n x_i (y_i - a - bx_i) = 0$$


From the first equation: $\sum y_i = na + b \sum x_i$, so $\hat{a} = \bar{y} - \hat{b}\bar{x}$

Substituting into the second equation and solving:

$$\hat{b} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

where $\bar{x} = \frac{1}{n} \sum x_i$ and $\bar{y} = \frac{1}{n} \sum y_i$.

The following visualizes this curve fitting approach:



This approach works as a pure optimization problem without probability.

2.5.2.3 Maximum likelihood solution

The probabilistic model provides the same result with additional benefits. Since $\epsilon_i \sim \text{Normal}(0, \sigma^2)$, we have:

$$y_i \mid x_i, a, b, \sigma^2 \sim \text{Normal}(a + bx_i, \sigma^2)$$

The likelihood for all observations is:

$$p(\mathbf{y} \mid \mathbf{x}, a, b, \sigma^2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(y_i - a - bx_i)^2}{2\sigma^2} \right\}$$

Taking the log-likelihood:

$$\log p(\mathbf{y} \mid \mathbf{x}, a, b, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - a - bx_i)^2$$

To maximize this with respect to a and b , we need to minimize:

$$\sum_{i=1}^n (y_i - a - bx_i)^2$$

This is exactly the same optimization problem as minimizing MSE!

Maximum likelihood estimation for linear regression with Normal errors is equivalent to ordinary least squares. The probabilistic model gives the same answer as curve-fitting but also provides a framework for quantifying uncertainty.

Therefore: $\hat{a} = \bar{y} - \hat{b}\bar{x}$ and $\hat{b} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$

These are identical to the formulas derived from minimizing MSE.

The following visualizes the maximum likelihood fit for the same data:

MLE Estimates:

Intercept: 1.92
Slope: 1.545
Sigma: 0.796

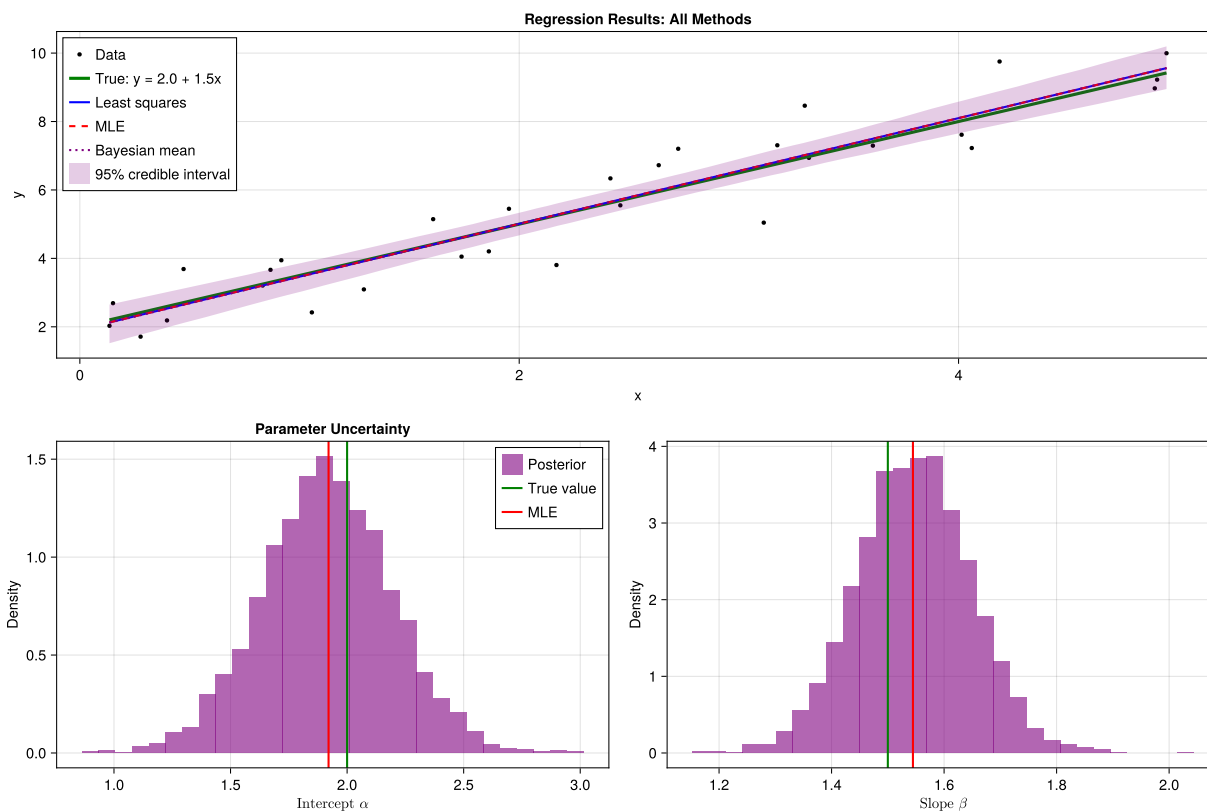
True Parameters:

Intercept: 2.0
Slope: 1.5
Sigma: 0.8

Least Squares Estimates (for comparison):

Intercept: 1.92
Slope: 1.545

A Bayesian fit with weakly informative priors yields similar point estimates with full posterior uncertainty; the following comparison summarizes results across approaches:



This connection between statistical and machine learning approaches extends far beyond linear regression. Many machine learning methods can be understood as maximum likelihood estimation under specific distributional assumptions, while statistical methods often reduce to familiar optimization problems. This curve-fitting perspective connects to the machine learning and non-parametric methods covered in [Chapter on Machine Learning](#), which address cases without specific functional form assumptions.

2.5.3 Logistic regression

2.5.3.1 Generative Story

The outcome y_i is a Bernoulli trial, where the probability of success p_i is related to a predictor x_i via the logistic function: $p_i = \frac{1}{1 + \exp\{-(\beta_0 + \beta_1 x_i)\}}$.

2.5.3.2 Maximum likelihood solution

No closed-form solution exists. Numerical optimization methods are required to find maximum likelihood estimates.

2.5.3.3 Bayesian solution

No conjugate priors exist for logistic regression. MCMC methods are required to sample from the posterior distribution.

2.5.3.4 Computational aspects

This example demonstrates problems where computational methods are essential rather than optional.

Further reading

For deeper study of probability and statistics:

- Blitzstein and Hwang (2019) provides excellent intuition with computational examples
- Downey (2021) emphasizes Bayesian thinking with practical applications
- Gelman (2021) connects regression to broader statistical modeling
- Gelman et al. (2014) comprehensive treatment of Bayesian computation
- [Computational Examples](#) demonstrates all methods with working code

3 Machine learning and nonparametric methods

Learning objectives

After reading this chapter, you should be able to:

- Apply fundamental supervised learning concepts to climate hazard assessment problems
- Understand nonparametric and semiparametric methods for flexible modeling
- Critically evaluate machine learning applications in climate risk literature
- Understand bias-variance tradeoffs and model validation approaches

3.1 Essential machine learning concepts

1. Supervised and unsupervised learning paradigms
2. Parametric vs nonparametric methods
3. Bias-variance tradeoff and regularization
4. Cross-validation and model selection
5. Tree-based methods and ensemble learning

Further reading

4 Correlation and dimensionality

Learning objectives

After reading this chapter, you should be able to:

- Model and interpret spatial dependence in climate fields
- Apply time series analysis methods to detect trends and patterns in climate data
- Use dimension reduction techniques for high-dimensional climate datasets
- Integrate spatial and temporal methods for spatiotemporal climate analysis

4.1 Essential concepts

1. Spatial statistics and geostatistical methods
2. Time series analysis and trend detection
3. Principal component analysis and empirical orthogonal functions
4. High-dimensional methods for climate data
5. Spatiotemporal integration approaches

Further reading

For spatial and temporal analysis in climate science:

- Cressie and Wikle (2011): Comprehensive treatment of spatial statistics

5 Model validation and comparison

Learning objectives

After reading this chapter, you should be able to:

- Apply graphical diagnostic methods to assess model fit quality
- Use information criteria (AIC, DIC, BIC) for quantitative model comparison
- Understand the relationship between model selection and predictive accuracy
- Recognize the subjective nature of model selection and make transparent choices

5.1 Essential model validation concepts

1. Graphical diagnostic methods and model checking
2. Information criteria for model comparison
3. Predictive accuracy and cross-validation
4. Model selection philosophy and transparency
5. Ensemble methods and model averaging

Further reading

- Piironen and Vehtari (2017): Technical treatment of predictive accuracy

Part III

II: Hazard Assessment

6 Extreme value theory

See first

This chapter builds on concepts from [Probability and Statistics](#).

Learning objectives


After reading this chapter, you should be able to:

- Understand the motivation for extreme value theory in climate risk assessment.
- Distinguish between block maxima and peak-over-threshold approaches.
- Select and fit appropriate extreme value distributions (GEV, GPD).
- Calculate return periods, return levels, and exceedance probabilities.
- Quantify and interpret different sources of uncertainty in extreme value analysis.
- Recognize the challenges posed by non-stationarity and climate change.
- Apply extreme value methods to real-world case studies.

6.1 Motivation

The design of reliable infrastructure, such as stormwater systems in Harris County, Texas, depends on a quantitative understanding of extreme environmental events. Engineers must characterize the magnitude and frequency of rare phenomena. For example, they might need to determine the expected recurrence interval of a daily rainfall total that exceeds a critical design threshold, such as 300 mm.

An initial, intuitive approach is to analyze the historical record.



Analysis of over 70 years of daily precipitation data from a station in Houston reveals 13 days where rainfall exceeded 150 mm. This allows for a simple empirical estimate of the event’s frequency. Thirteen exceedances in 71 years suggests an average recurrence interval of approximately 5.5 years. For the 300 mm threshold, however, the historical record contains zero events. An empirical estimate of the probability is therefore zero, which is uninformative for risk assessment and infrastructure design. We require a method for extrapolating beyond the range of observed data.

A natural extension would be to fit a standard parametric probability distribution—such as a Gamma distribution, which is commonly used for precipitation—to the entire set of daily rainfall observations. One could then use the extreme upper tail of the fitted distribution to estimate the probability of exceeding 300 mm.

This approach, however, is fundamentally unreliable. The goodness-of-fit of a parametric model is driven by its ability to describe the bulk of the data, where observations are plentiful. Minor model misspecifications in this high-density region can translate into substantial errors in the far tails of the distribution, where data are sparse or nonexistent. As stated by Coles (2001), the foundational text on this topic, “very small discrepancies in the estimate of F can lead to substantial discrepancies for F^n ”, where in our example F is the cumulative distribution function (CDF) of daily rainfall while F^n is the CDF of the maximum daily rainfall over n days.

The models that best describe the central tendency of a process are not necessarily suitable for describing its extremes. This fragility of tail extrapolation necessitates a theoretical framework developed specifically for modeling the behavior of extreme values.

A naive empirical estimator for the exceedance probability of a value x based on n observations X_i is the simple proportion of exceedances:

$$\hat{P}(X > x) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(X_i > x)$$

where $\mathbf{1}(\cdot)$ is the indicator function. This is useful when the number of exceedances is large, but fails for extrapolation when the sum is zero, as was the case for the 300mm threshold. This mathematically demonstrates the limitations of simple empirical methods and motivates the specialized approaches that follow.

6.2 Approaches for modeling extreme values

Extreme Value Theory (EVT) provides such a framework. Rather than modeling the entire distribution of a process, EVT focuses on the distribution of its extreme values. There are two primary methods for extracting these values from a time series.

6.2.1 Block maxima

In this method, the period of observation is divided into non-overlapping blocks of equal duration (e.g., years). For each block of size n , we define the block maximum as $M_n = \max\{X_1, \dots, X_n\}$. This yields a new time series of block maxima. A crucial conceptual point is that we are not studying a single maximum possible value of the process. Instead, we are studying the distribution of the sample maximum, M_n , which is itself a random variable. This is why a distribution like the GEV is needed to describe its behavior. This approach is intuitive and directly relates to concepts of annual risk, but it can be inefficient as it discards other potentially extreme events within a block.

6.2.2 Peaks-over-threshold

The peaks-over-threshold (POT) method defines a set of exceedances over a high threshold u as $\{X_i : X_i > u\}$. The variable of interest is the value of the **excesses** themselves, $Y = X_i - u$, for all X_i in the set of exceedances. The GPD is a model for these excess amounts, not the raw values. This approach is more data-efficient than the block maxima method. Its primary challenge lies in selecting an appropriate threshold, which involves a trade-off between bias (if the threshold is too low) and variance (if the threshold is too high, yielding too few data points). In practice, exceedances may occur in clusters (e.g., during a multi-day storm event). Therefore, a declustering algorithm is often applied to the raw exceedances to ensure that the events being modeled are approximately independent.

6.3 Asymptotic theory for extremes

The statistical justification for both the block maxima and POT approaches comes from asymptotic theorems that describe the limiting distributions of extreme values.

6.3.1 Theory for block maxima: The GEV distribution

The theoretical justification for the block maxima approach is the **Extremal Types Theorem**. This theorem states that for a large class of underlying distributions, if a stable, non-degenerate limiting distribution for the block maximum M_n exists, it must be the **Generalized Extreme Value (GEV)** distribution.

A critical point is that the theorem applies to a *linearly rescaled* or *normalized* maximum. Consider the raw maximum, M_n . As the block size n increases, the expected value of M_n will also increase (or stay the same)—its distribution drifts towards larger values and does not converge. A similar issue arises in the Central Limit Theorem (CLT), which describes the convergence of a sample mean. The CLT does not apply to the raw sum of random variables, but to a normalized version of it.

The Extremal Types Theorem is the direct analog of the CLT for maxima. It states that there exist sequences of normalizing constants, a location parameter b_n and a scale parameter $a_n > 0$, such that the distribution of the normalized maximum, $(M_n - b_n)/a_n$, converges to the GEV distribution as $n \rightarrow \infty$.

In practice, we do not need to know the specific functional forms of a_n and b_n . Instead, for a fixed, large block size n (e.g., one year of daily data), we fit the GEV distribution directly to the series of block maxima. The GEV's location and scale parameters, μ and σ , effectively absorb and account for the normalization that the theory requires. The GEV is a flexible three-parameter family with a cumulative distribution function (CDF) given by:

$$G(z) = \exp \left\{ - \left[1 + \xi \left(\frac{z - \mu}{\sigma} \right) \right]^{-1/\xi} \right\}$$

This is defined on the set $\{z : 1 + \xi(z - \mu)/\sigma > 0\}$. The parameters are location (μ), scale ($\sigma > 0$), and shape (ξ). These parameters implicitly depend on the block size n .

6.3.2 Theory for threshold exceedances: The GPD

The corresponding theory for the POT approach is the **Pickands–Balkema–de Haan Theorem**. It states that for a wide range of distributions, the distribution of excesses over a high threshold u can be approximated by the **Generalized Pareto Distribution (GPD)**. Conceptually, if the GEV describes the behavior of the maximum of a large block, the GPD describes the behavior of the distribution's tail that produced that maximum. It is the distribution one would expect to see by “zooming in” on the tail of a distribution above a high threshold.

The GPD is a two-parameter family with a CDF for **excesses** $Y = X - u$ given by:

$$H(y) = 1 - \left(1 + \frac{\xi y}{\sigma_u} \right)^{-1/\xi}$$

This is defined on $\{y : y > 0 \text{ and } (1 + \xi y/\sigma_u) > 0\}$. The parameters are the shape, ξ , and a scale parameter σ_u that depends on the threshold u .

If the parent distribution's maxima are GEV-distributed with parameters (μ, σ, ξ) , then the GPD scale parameter is given by $\sigma_u = \sigma + \xi(u - \mu)$. This key result shows that the GPD scale parameter is a linear function of the threshold u , a property that is used in advanced diagnostics for threshold selection, and illustrates the links between the GPD and GEV models for extremes.

6.3.3 The shape parameter

The **shape parameter**, ξ , is the most critical parameter in extreme value analysis and has the same interpretation in both the GEV and GPD. It governs the tail behavior of the distribution.

- $\xi = 0$ (**Gumbel-type tail**): The tail decays exponentially (light tail).
- $\xi > 0$ (**Fréchet-type tail**): The tail decays as a polynomial (heavy tail), with no upper bound.
- $\xi < 0$ (**Weibull-type tail**): The distribution has a finite upper bound.

6.3.4 Connection between GEV and GPD

The GEV and GPD models are intrinsically linked. If the block maxima of a process follow a GEV distribution with parameters (μ, σ, ξ) , then for a high threshold u , the threshold excesses follow a GPD with the same shape parameter ξ . This provides a theoretical consistency between the two primary modeling approaches.

6.4 Return periods and return levels

6.4.1 Definitions and Calculation

The output of an extreme value analysis is typically expressed in terms of **return periods** and **return levels**. The **N-year return level**, z_N , is the level expected to be exceeded on average once every N years. It corresponds to the quantile of the distribution with an annual exceedance probability of $p = 1/N$.

For the GEV distribution, it is calculated by inverting the CDF:

$$z_N = \mu - \frac{\sigma}{\xi} [1 - (-\ln(1 - p))^{-\xi}]$$

Calculating return levels from a GPD model requires an additional step. If ζ_u is the rate of threshold exceedances (e.g., the average number of exceedances per year), the N -year return level is the value z_N that is exceeded with probability $1/(N\zeta_u)$. It is calculated by adding the threshold back to the corresponding quantile of the GPD:

$$z_N = u + \frac{\sigma_u}{\xi} [(N\zeta_u)^\xi - 1]$$

6.4.2 Return level plots

A standard diagnostic and visualization tool is the return level plot. This plot graphs the estimated return level z_N against the return period N , with N typically plotted on a logarithmic scale. The curvature of the fitted line is a direct visualization of the shape parameter, ξ : a straight line implies $\xi = 0$, a concave curve implies $\xi > 0$, and a convex curve implies $\xi < 0$.

[Placeholder for a Return Level Plot showing a fitted GEV or GPD curve with confidence intervals.]

6.5 Inference

6.5.1 Plotting Positions

To visually assess model fit, the fitted model is plotted against the observed maxima. For this purpose, we require a method to assign a non-exceedance probability (and thus a return period) to each observed maximum. For a sample of n block maxima, let $z_{(1)} < z_{(2)} < \dots < z_{(n)}$ be the ordered values. We estimate the probability $P(Z \leq z_{(k)})$ using a plotting position formula. These formulas generally take the form:

$$p_k = \frac{k - a}{n + 1 - 2a}$$

where k is the rank of the observation and a is a parameter. Common choices include:

- **Weibull:** $a = 0$, giving $p_k = k/(n + 1)$.
- **Gringorten:** $a = 0.44$, giving $p_k = (k - 0.44)/(n + 0.12)$. This is often recommended as an unbiased choice for Gumbel-type distributions.

The empirical return period for the k -th observation is then estimated as $1/(1 - p_k)$.

6.5.2 Moments

6.5.3 MLE

6.5.4 Bayesian

6.6 Sampling variability

6.7 Regionalization

6.8 Nonstationarity

Further reading

- (Coles 2001): Canonical extreme value textbook with mathematical rigor and practical examples

7 Downscaling and Bias Correction

Learning objectives

After reading this chapter, you should be able to:

- Distinguish between supervised and distributional downscaling approaches
- Understand the motivation for downscaling climate model outputs
- Apply bias correction and quantile-quantile mapping techniques
- Recognize the stationarity assumption and its implications
- Evaluate different downscaling methods for specific applications
- Understand modern machine learning approaches to climate downscaling

Further reading

- Lanzante et al. (2018) for comprehensive review of downscaling challenges
- Lafferty and Sriver (2023)
- Farnham, Doss-Gollin, and Lall (2018)

8 Stochastic weather generators

Learning objectives

After reading this chapter, you should be able to:

- Understand the principles of stochastic weather generation
- Apply statistical models for synthetic weather data
- Use weather generators for downscaling climate model output
- Evaluate the performance of weather generation models

8.1 Essential concepts

1. Hidden Markov models for weather state modeling
2. Multi-site and multi-variable generation
3. Statistical downscaling applications
4. Synthetic data validation and evaluation
5. Integration with physical models

Further reading

9 Physics-based models and calibration

Learning objectives

After reading this chapter, you should be able to:

- Navigate trade-offs between model complexity, interpretability, and computational cost
- Characterize and communicate within- and between-model uncertainty
- Use surrogate models to approximate complex model output
- Apply model calibration techniques for climate applications

9.1 Essential concepts

1. Physics-based vs data-driven modeling spectrum
2. Model chaining and uncertainty propagation
3. Calibration methods and parameter estimation
4. Surrogate modeling for computational efficiency
5. Model structure uncertainty quantification

Further reading

For physics-based modeling in climate applications:

- Rackauckas et al. (2020): Scientific machine learning for differential equations

10 Optimal sampling methods

Learning objectives

After reading this chapter, you should be able to:

- Apply sampling techniques to generate synthetic event sets
- Use importance and stratified sampling to improve efficiency in hazard modeling
- Evaluate how sampling choices affect estimates of extreme risk
- Balance computational cost vs accuracy in climate risk estimation

10.1 Essential sampling concepts

1. Monte Carlo sampling and variance reduction
2. Importance sampling for rare events
3. Stratified sampling strategies
4. Synthetic event generation methods
5. Computational efficiency optimization

Further reading

11 Global sensitivity analysis

Learning objectives

After reading this chapter, you should be able to:

- Understand the role of sensitivity analysis in climate risk modeling
- Apply variance-based sensitivity methods (Sobol indices) to identify key parameters
- Use Morris screening methods for initial parameter importance ranking
- Interpret sensitivity analysis results for model simplification and uncertainty reduction

11.1 Essential sensitivity analysis concepts

1. Global vs local sensitivity analysis approaches
2. Sobol indices and variance decomposition
3. Morris elementary effects for screening
4. Multi-model sensitivity analysis
5. Implementation strategies for complex model chains

Further reading

For global sensitivity analysis:

- Saltelli et al. (2008): Comprehensive introduction to GSA methods
- Herman and Usher (2017): Practical implementation guide with software tools

Part IV

III: Risk Management

12 Exposure and Vulnerability

12.1 Learning objectives

By the end of this chapter, you should be able to:

- Define exposure and vulnerability in the context of climate risk assessment
- Distinguish between different types of vulnerability (physical, social, economic)
- Understand methods for quantifying and mapping exposure
- Apply vulnerability assessment frameworks to real-world scenarios
- Integrate exposure and vulnerability data with hazard information for risk assessment

12.2 Further reading

13 Cost-Benefit Analysis and Net Present Value

Learning objectives

After reading this chapter, you should be able to:

- Understand the theoretical foundation of cost-benefit analysis and Bayesian decision theory
- Apply net present value calculations with appropriate discount rates
- Quantify costs and benefits using utility functions for climate risk decisions
- Handle uncertainty in cost-benefit frameworks using expected value
- Recognize the limitations and appropriate applications of cost-benefit analysis
- Evaluate economic trade-offs over different time horizons and scenarios

Further reading

14 Policy Search & Optimization

15 Risk Transfer

See first

This chapter builds on concepts from: - [Exposure and Vulnerability](#) - [Expectations and Discounting](#)

Learning objectives

- Explore financial instruments (insurance, reinsurance, catastrophe bonds) for spreading or transferring climate risk.
- Understand parametric insurance triggers and how they differ from indemnity-based approaches.
- Assess the role of public-private partnerships in risk transfer mechanisms.

15.1 Insurance/reinsurance fundamentals

15.2 Parametric vs. indemnity coverage

15.3 Catastrophe bonds, index-based schemes

15.4 Challenges in emerging markets and vulnerable regions

Further reading

16 Deep Uncertainty and Model Structure

See first

This chapter builds on concepts from: - [Probability and Statistics](#) - [Model Validation and Comparison](#)

Learning objectives

- Distinguish between aleatory and epistemic uncertainty in climate risk assessment
- Understand the challenges posed by structural uncertainty and model disagreement
- Apply Bayesian Model Averaging (BMA) and stacking approaches to combine multiple models
- Recognize when deep uncertainty invalidates traditional decision frameworks
- Identify sources of deep uncertainty in exposure and impact modeling

Further reading

17 Robustness

18 Adaptive Planning and Flexibility

See first

This chapter builds on concepts from: - [Deep Uncertainty](#) - [Robustness](#)

Learning objectives

- Plan for uncertainty with adaptive management and iterative risk strategies.
- Develop adaptation pathways that evolve with new information (e.g., climate data, impacts).
- Incorporate monitoring and feedback loops into long-term climate policy.

Further reading

19 Working with People: Values, Participation, and Communication

Learning objectives

- Understand and communicate how decision analyses are necessarily influenced by subjective value judgements
- Explain common ethical frameworks for decision analysis
- Be familiar with multiple frameworks for stakeholder engagement and participatory methods
- Design effective approaches for communicating uncertainty in climate risk assessments
- Recognize the social and political dimensions of climate risk management

Further reading

Part V

Computational Case Studies

Overview

This collection of computational notebooks demonstrates the methods and concepts discussed in the main text through practical applications. Each notebook is designed to be standalone and self-contained, using the Julia programming language for all computations.

Probability, inference, and computation: examples

This notebook provides computational examples to accompany the [Probability and inference](#) chapter. See also: [Moquito bites and beer consumption](#)

1. We use `using` to import necessary packages

Distribution examples


Understanding probability distributions requires grasping three fundamental functions: the probability density function (PDF) or probability mass function (PMF), the cumulative distribution function (CDF), and the quantile function. These examples demonstrate how these functions work together using the normal distribution (continuous) and Poisson distribution (discrete).

The visualizations show both forward operations (finding probabilities from values using the CDF) and inverse operations (finding values from probabilities using quantiles). This computational perspective reinforces the theoretical relationships discussed in the main chapter.

We start by defining some helper functions.

1. Creates function to shade area under PDF curve between bounds a and b . By convention a function that ends in `!` modifies its arguments (in this case, a plot axis).
2. Vectorized evaluation of PDF using broadcasting (the `.` operator)
3. Calculate probability using CDF difference: $P(a \leq X \leq b) = F(b) - F(a)$
4. Forward CDF operation: given x , find $F(x) = P(X \leq x)$
5. Inverse CDF operation: given probability p , find x such that $F(x) = p$
6. The quantile function is the inverse of the CDF

Next, we use [Makie.jl](#) to creat the plot.



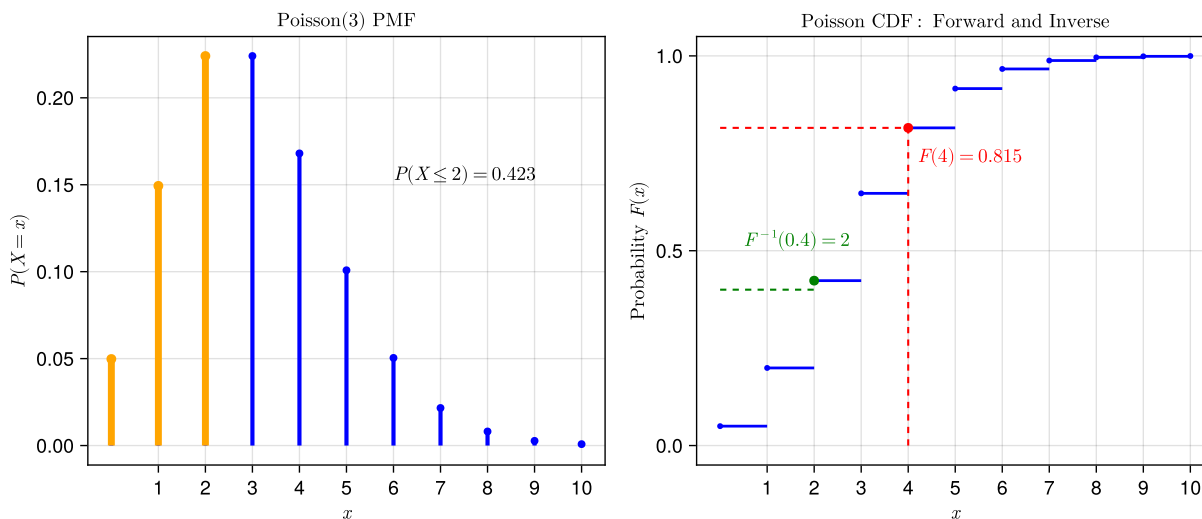
1. Create `Normal(0,1)` distribution object from `Distributions.jl`
2. Plot PDF using vectorized evaluation over `x_range`
3. Plot CDF showing cumulative probabilities

The normal distribution demonstrates these concepts for continuous variables, where probabilities correspond to areas under smooth curves. Next, we examine the Poisson distribution for discrete random variables, where probabilities correspond to point masses and the CDF becomes a step function.

We define some more helper functions for discrete distributions.

1. Plots discrete PMF as stems (vertical lines) with points at the top
2. For discrete distributions, `pdf()` returns probability mass $P(X = x)$
3. Draw vertical line from 0 to probability mass at each x value
4. Highlights specific probability masses in different color
5. Sum individual masses to get cumulative probability $P(X \leq 2)$
6. Creates step function visualization for discrete CDF
7. Horizontal line segments show constant CDF values between integers
8. Same forward operation logic but adapted for discrete distributions
9. Inverse operation finds the smallest integer x such that $F(x) \geq p$

Now we create the plot for the Poisson distribution.



1. Create Poisson distribution with rate parameter $\lambda = 3$
2. Calculate and highlight $P(X \leq 2)$ by summing individual probability masses


Coin flip examples

Having explored the foundational concepts of probability distributions, we now turn to statistical inference—the process of learning about unknown parameters from observed data.

The coin flip example provides a pedagogically clean introduction to statistical inference. We observe a series of coin flips and want to infer the probability θ of heads. This simple setup allows us to demonstrate maximum likelihood estimation, Bayesian inference, and the relationship between analytical and computational approaches.

Likelihood visualization

The likelihood function shows how plausible different values of θ are given our observed data. For the Binomial model, this creates a clear peak at the observed proportion of heads.

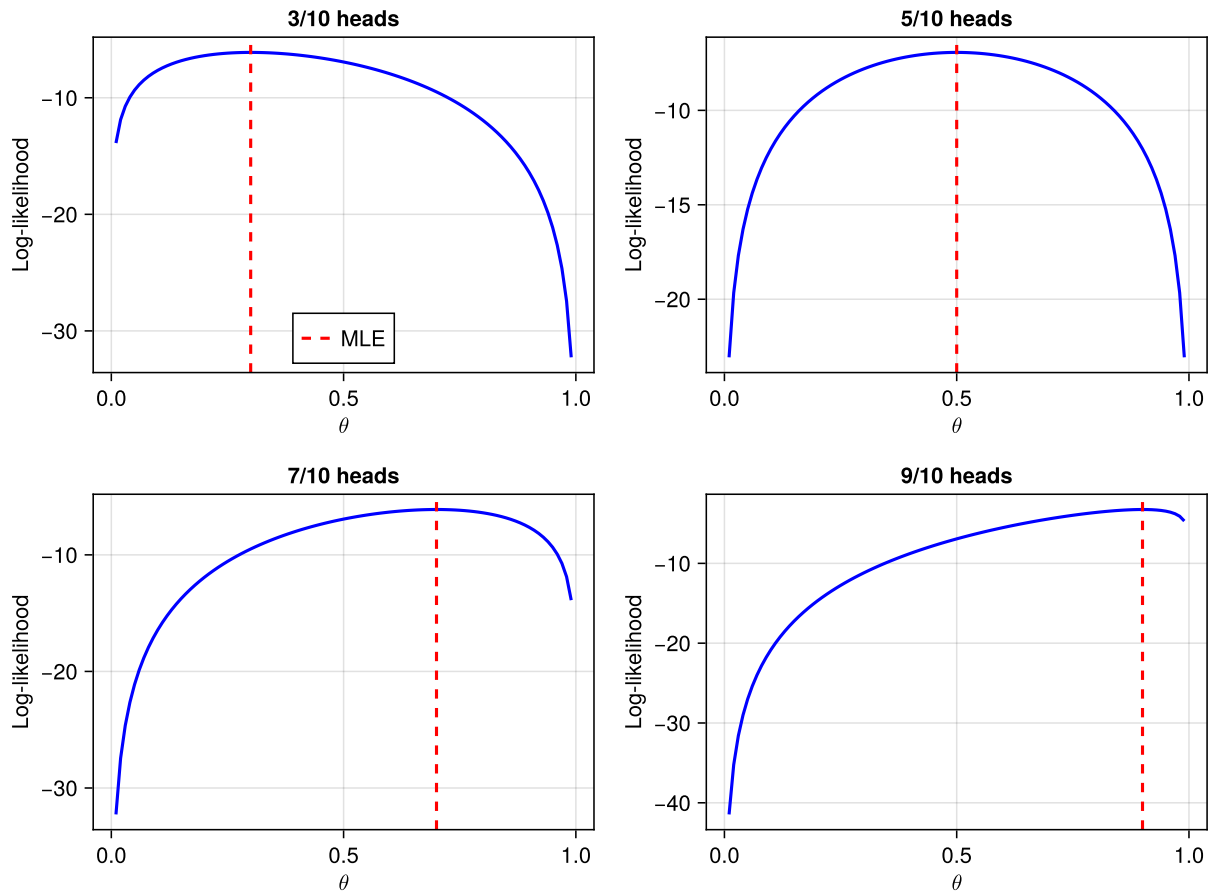


1. Create grid of θ values from 0.01 to 0.99 to evaluate likelihood
2. Evaluate binomial likelihood $L(\theta) = P(y|\theta, n)$ for each θ value
3. Maximum likelihood estimate is simply the observed proportion $\frac{y}{n}$

The likelihood peaks at $\theta = 0.7$ when we observe 7 heads out of 10 flips. This demonstrates that the maximum likelihood estimate is simply the observed proportion.

Maximum likelihood estimation

Maximum likelihood estimation finds the parameter value that makes the observed data most probable. For the coin flip problem, we can solve this analytically by taking the derivative of the log-likelihood and setting it to zero. The following shows how the MLE varies with different numbers of observed heads:



1. Loop through different numbers of observed heads to show MLE behavior
2. Create 2×2 grid of subplots using integer division and modulo
3. Analytical log-likelihood: $\ell(\theta) = y \ln(\theta) + (n - y) \ln(1 - \theta)$
4. Demonstrate MLE for 3, 5, 7, and 9 heads out of 10 flips

Each plot shows the log-likelihood function as a smooth curve with its maximum clearly indicated. Notice that the MLE is always at the observed proportion $\frac{y}{n}$, confirming our analytical derivation.

Bayesian analysis

Bayesian inference treats the parameter θ as a random variable with its own probability distribution. We start with a prior belief about θ , observe data, and update our beliefs to get the posterior distribution.

For the coin flip problem with a **Beta** prior, we can compute the posterior analytically using conjugacy. The posterior parameters are simply the prior parameters plus the observed successes and failures. This example compares the exact analytical solution with MCMC sampling.

First we define the `Turing.jl` model and the comparison function. See the official [Turing.jl documentation and tutorials](#).

coin_flip_model (generic function with 2 methods)

1. Turing.jl probabilistic programming model specification using @model macro
2. Prior distribution: $\theta \sim \text{Beta}(\alpha, \beta)$
3. Likelihood: $y \sim \text{Binomial}(n, \theta)$ connects data to parameter

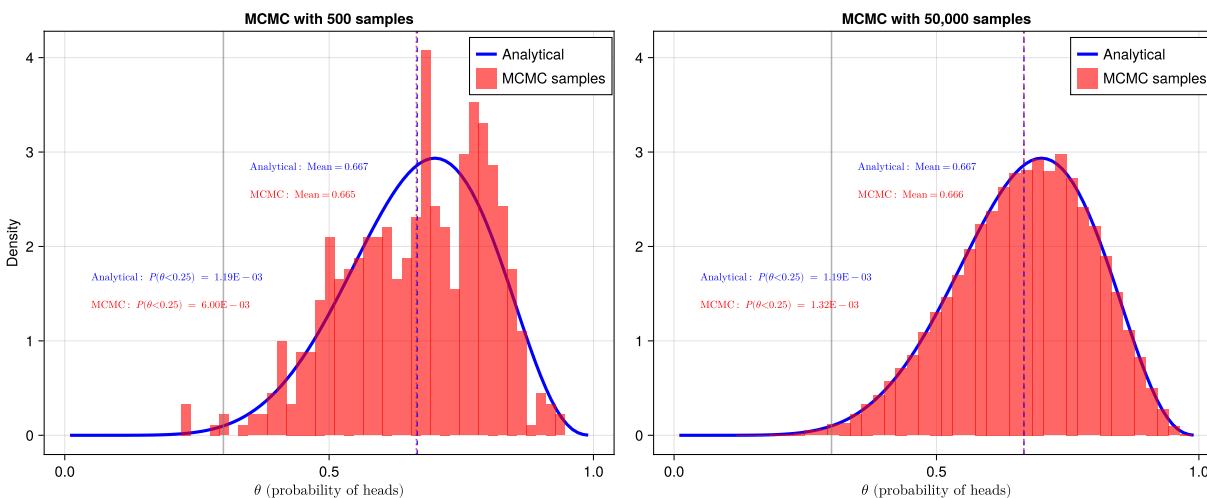
Next we define a function to help us create the data and fits

```
Info: Found initial step size
      = 1.6
Info: Found initial step size
      = 1.6
```

```
(analytical_posterior = Distributions.Beta{Float64}(=8.0, =4.0), mcmc_samples = [0.8072931234
```

4. Function comparing analytical vs MCMC approaches
5. Posterior $\alpha = \text{prior } \alpha + \text{successes}$ (conjugate update)
6. Posterior $\beta = \text{prior } \beta + \text{failures}$ (conjugate update)
7. Exact posterior using Beta-Binomial conjugacy
8. Create Turing model instance with data and priors
9. MCMC sampling using No-U-Turn Sampler (NUTS)
10. Extract θ samples from MCMC chain using `Array(chain[:])`
11. Low sample count (500) to show MCMC variability
12. High sample count (50,000) to show MCMC convergence

And then we plot



13. Link axes for direct comparison between panels
14. Loop over both low and high sample datasets
15. Exact analytical posterior density
16. Normalized histogram of MCMC samples
17. Exact tail probability using CDF
18. Format tail probability in scientific notation

19. Estimate tail probability from MCMC samples

The results show excellent agreement between analytical and MCMC approaches. With sufficient samples, MCMC accurately approximates the true posterior distribution. The comparison also illustrates how MCMC precision varies across the distribution—central quantities like the mean are estimated more accurately than tail probabilities.

Linear regression example

The coin flip example demonstrated inference for a single parameter using a simple likelihood function. Real-world problems typically involve multiple parameters and more complex relationships between variables. Linear regression provides our first example of multivariate statistical inference while maintaining analytical tractability.


This example demonstrates the progression from curve fitting to probabilistic inference. We'll examine the same problem through three lenses: minimizing prediction error, maximum likelihood estimation, and Bayesian inference. All three approaches yield the same parameter estimates for the linear model with normal errors, but provide different ways of quantifying uncertainty.

Generating synthetic data

We start by creating synthetic data from a known linear relationship with added noise. This allows us to verify that our inference methods recover the true parameters.

1. Set random seed for reproducible results
2. Generate sorted predictor values from uniform distribution
3. Generate response following $y_i = \alpha + \beta x_i + \epsilon_i$ where $\epsilon_i \sim N(0, \sigma^2)$


Let's visualize our synthetic dataset:



This visualization shows our synthetic dataset with the true linear relationship and noise level. The scattered points represent our “observations,” while the red line shows the true underlying relationship we’re trying to recover. The shaded band indicates the expected range of observations given the noise level ($\pm 2\sigma$).

Curve fitting by minimizing prediction error

Without probability theory, we can view regression as finding the line that best fits the data. This requires choosing a loss function to measure fit quality. The most common choice is mean squared error (MSE), which penalizes large prediction errors more than small ones. This is more of a “machine learning” perspective on the problem, and is one to which we return in [the ML chapter](#).



1. Calculate least squares estimates using analytical formulas
2. Slope: $\hat{\beta} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$
3. Intercept: $\hat{\alpha} = \bar{y} - \hat{\beta}\bar{x}$ (fitted line passes through centroid)
4. Mean squared error quantifies overall fit quality
5. Residual lines show prediction errors for each observation

The least squares method finds parameter estimates close to the true values. The residual lines visualize the model's prediction errors, with shorter lines indicating better fit.

Maximum likelihood estimation

The curve fitting approach works but provides no way to quantify uncertainty in our parameter estimates. By introducing a probabilistic model, we can use maximum likelihood estimation to both estimate parameters and assess their precision.

For linear regression, we assume the response follows a normal distribution around the linear predictor:

$$y_i \sim N(\alpha + \beta x_i, \sigma^2)$$

as discussed in the theory section. Using the **Optim** package, we can numerically maximize the likelihood function to find the MLE estimates for α , β , and σ . In **Optim**, all functions are minimized, so we minimize the negative log-likelihood instead.

MLE Estimates:

Intercept: 1.92

Slope: 1.545

Sigma: 0.796

True Parameters:

Intercept: 2.0

Slope: 1.5

Sigma: 0.8

Least Squares Estimates (for comparison):

Intercept: 1.92

Slope: 1.545

1. Negative log-likelihood function for optimization (we minimize this)
2. Enforce constraint that $\sigma > 0$ by returning infinite cost for invalid values
3. Sum negative log-likelihoods: $-\sum_i \log p(y_i | \alpha + \beta x_i, \sigma)$
4. Function to find MLE using numerical optimization
5. Use Nelder-Mead simplex algorithm for unconstrained optimization

The MLE estimates for the intercept and slope match the least squares results exactly. This confirms the theoretical result that for linear regression with normal errors, maximum likelihood estimation is equivalent to minimizing mean squared error. The MLE approach additionally estimates σ , quantifying the noise level in the data.

Bayesian inference

Bayesian inference treats all parameters as random variables with probability distributions. We specify prior beliefs, observe data, and compute posterior distributions that quantify parameter uncertainty.

Info: Found initial step size
= 0.0125

Bayesian Posterior Summaries:

Intercept: 1.912 \pm 0.28

Slope: 1.546 \pm 0.101

Sigma: 0.846 \pm 0.12

1. Turing.jl probabilistic programming model using `@model` macro
2. Weakly informative prior for intercept: $\alpha \sim N(0, 5^2)$
3. Weakly informative prior for slope: $\beta \sim N(0, 5^2)$
4. Exponential prior for noise standard deviation: $\sigma \sim \text{Exp}(1)$
5. Likelihood specification: each $y_i \sim N(\alpha + \beta x_i, \sigma)$
6. Sample from posterior using No-U-Turn Sampler (NUTS)
7. Extract parameter samples from MCMC chain and flatten to vectors

MLE with Turing

We can also use `Turing.jl` to find the MLE by optimizing the model instead of sampling. This demonstrates how the same probabilistic model can be used for both optimization-based and sampling-based inference.

Turing MLE Estimates:

Intercept: 1.916

Slope: 1.546

Sigma: 0.785

Comparison with Manual MLE:

Intercept difference: 0.004083

Slope difference: 0.001082

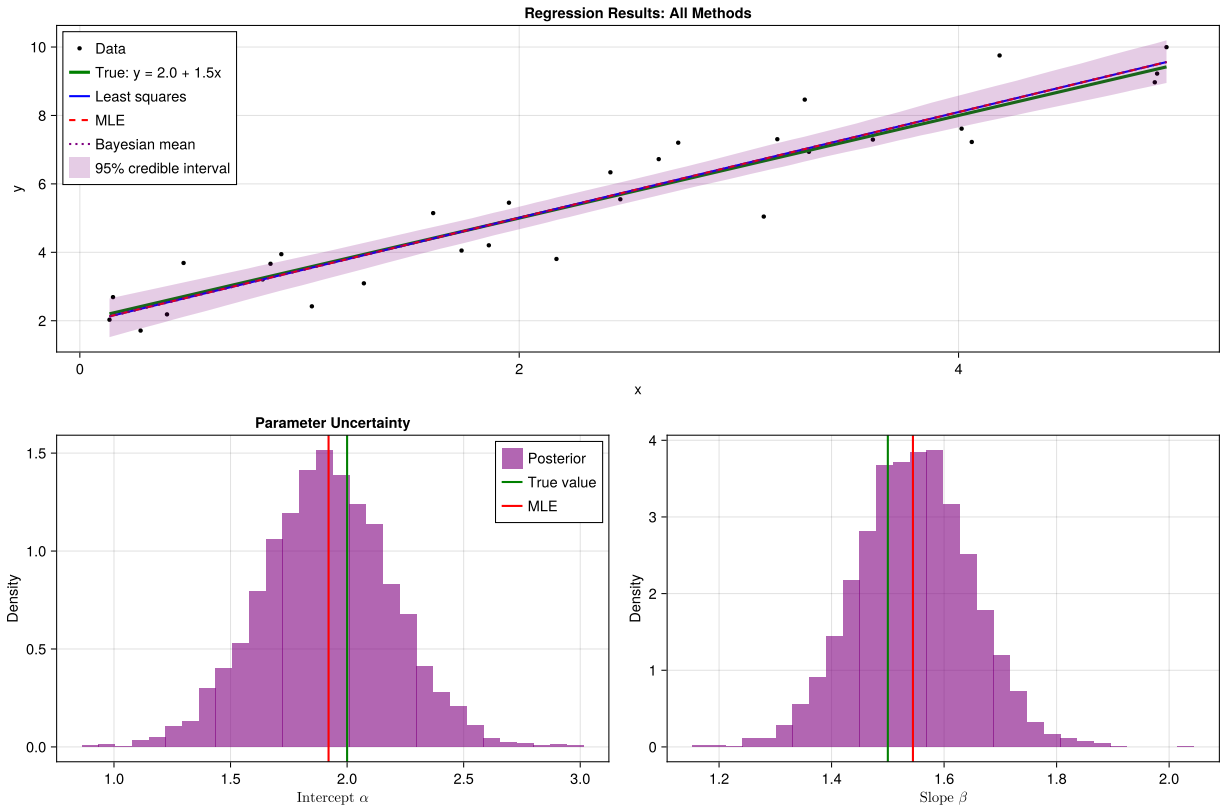
Sigma difference: 0.010205

1. Use the same probabilistic model defined earlier for Bayesian inference
2. Optimize using `MAP()` (Maximum A Posteriori) which reduces to MLE with flat priors

This approach shows that `Turing.jl` provides a unified framework: the same model definition can be used for both MLE optimization and MCMC sampling. The results should match our manual MLE implementation, demonstrating the consistency between different computational approaches, although small numerical differences may arise due to optimization tolerances.

Comparison of approaches

All three approaches provide similar point estimates, but differ in their treatment of uncertainty:



1. Use subset of MCMC samples for computational efficiency
2. For each x-value, calculate predicted y using different parameter samples
3. Calculate 2.5th and 97.5th percentiles for 95% credible interval

This comparison reveals several important points:

- **Point estimates agree:** All methods produce nearly identical parameter estimates
- **Uncertainty quantification differs:** Only the Bayesian approach provides full uncertainty distributions and credible intervals
- **Prediction uncertainty:** The 95% credible interval shows the range of plausible regression lines given parameter uncertainty
- **Computational cost varies:** Least squares is fastest, MLE requires optimization, MCMC is most expensive
- **Interpretability:** Bayesian posteriors directly quantify “how certain should we be about these parameter values?”


The Bayesian approach excels when we need to propagate parameter uncertainty through subsequent calculations or decision processes. The credible interval demonstrates how parameter uncertainty translates into prediction uncertainty—a crucial consideration for risk assessment and decision-making under uncertainty.

Extreme Value Theory Examples

This notebook demonstrates extreme value analysis concepts using real Houston precipitation data, illustrating the fundamental extrapolation problem in climate risk assessment.

Shape parameter implications

The shape parameter ξ in extreme value distributions fundamentally determines how heavy the tail is, which directly affects our extrapolation to rare events. Let's visualize how different shape parameter values lead to dramatically different tail behavior:



The key insight: the shape parameter ξ controls how rapidly probabilities decrease in the tail.

Houston precipitation data

Let's look at this using a realistic example

Loaded 70366 raw GHCN daily observations

First few rows:

	USC00414313	19430101	PRCP	0	Column5	Column6	6	Column8
	String15	Int64	String7	Int64	String1?	Missing	String1	Int64?
1	USC00414313	19430102	PRCP	0	<i>missing</i>	<i>missing</i>	6	<i>missing</i>
2	USC00414313	19430103	PRCP	0	<i>missing</i>	<i>missing</i>	6	<i>missing</i>
3	USC00414313	19430104	PRCP	0	<i>missing</i>	<i>missing</i>	6	<i>missing</i>

This loads the raw GHCN (Global Historical Climatology Network) data file. Each row contains weather measurements with a specific format that we need to decode.

Processing GHCN format

GHCN daily files use a fixed format where each row represents one measurement at one station on one date:

After filtering for precipitation (PRCP):

- Total PRCP observations: 25225
- Date range (raw): 19430102 to 20131130

GHCN includes many weather variables (temperature, snow, etc.), but we only need precipitation (PRCP). The date is stored as YYYYMMDD format and values are in tenths of millimeters.

Converting dates and units

After format conversion:

- Daily observations: 25225
- Date range: 1943-01-02 to 2013-11-30
- Years of data: 71
- Max daily rainfall: 272.3 mm

Now we have properly formatted dates and rainfall measurements in standard millimeter units. This represents over 70 years of daily precipitation measurements - our "high-frequency observations."

Filling missing dates for complete time series

GHCN data only includes dates where measurements were recorded. For proper time series analysis and plotting, we need to fill in missing dates:


Date completeness:

- GHCN observations: 25225 days
- Complete date range: 25901 days
- Missing days: 676
- Complete time series: 25901 days
- Missing rainfall values: 676


This creates a complete daily time series from $\$(\text{minimum}(\text{precip_data.date}))$ to $\$(\text{maximum}(\text{precip_data.date}))$ with `missing` values for days without measurements. Missing values are common in historical weather data due to equipment issues, observer absence, or data quality problems.

Daily precipitation time series

Now let's visualize the daily precipitation data to see the “high-frequency observations” mentioned in our extrapolation problem:



Or a subset



Statistics Without the Agonizing Details

Learning Objectives

- Compare simulation-based and parametric statistical tests
- Understand the logic of permutation testing
- Apply computational methods to hypothesis testing
- Interpret p-values through simulation

In this class we will use computation and simulation to build fundamental insight into statistical processes without dwelling on “agonizing” details. Here we implement the excellent example problem from John Rauser. First, watch the video:

We will recreate his analysis to answer the fundamental question:

Does drinking beer reduce the likelihood of being bitten by mosquitos?

Data

First, we will create the data. Here is the data for the beer drinkers:

Using Julia, we can learn more about the data through exploratory analysis.

23.6

We can do the same for water drinkers:

A simple analysis

Following Rauser, let’s calculate the difference between the average number of bites in each group.

The skeptic's argument

The skeptic asks whether this might be random chance.

1. We could answer this with a T test:

- Determine if there is a significant difference between the means of two groups
- Assumes (approximate) normality
- Assumptions hidden behind a software package

2. Simulation approach:

- Suppose the skeptic is right – the two groups are sampled from the same distribution
- Shuffle the data (randomly divide into two groups by assuming that there is no difference between the two groups)
- Calculate the difference between each group
- Repeat many times and examine the distribution of differences

Implementation


We can code up the simulation approach in Julia.

We want to learn about the *sampling distribution* of the group differences: repeat this experiment many times over and plot the results.

50000

Plotting

```
Warning: Found `resolution` in the theme when creating a `Scene`. The `resolution` keyword fo
@ Makie ~/.julia/packages/Makie/Q6F2P/src/scenes.jl:238
```



Alternative

We could have done this with a parametric test

0.00056

References

- Abernathey, Ryan. 2024. *An Introduction to Earth and Environmental Data Science*. <https://earth-env-data-science.github.io/intro.html>.
- Applegate, Patrick, and Klaus Keller. 2015. *Risk Analysis in the Earth Sciences*. Leanpub. <https://leanpub.next/raes>.
- Bastani, Hamsa, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakcı, and Rei Mariman. 2025. “Generative AI Without Guardrails Can Harm Learning: Evidence from High School Mathematics.” *Proceedings of the National Academy of Sciences* 122 (26): e2422633122. <https://doi.org/10.1073/pnas.2422633122>.
- Bishop, Christopher M., and Hugh Bishop. 2024. *Deep Learning: Foundations and Concepts*. Cham: Springer International Publishing. <https://doi.org/10.1007/978-3-031-45468-4>.
- Blitzstein, Joseph K., and Jessica Hwang. 2019. *Introduction to Probability, Second Edition*. 2nd Edition. Boca Raton: Chapman and Hall/CRC. <http://probabilitybook.net>.
- Coles, Stuart. 2001. *An Introduction to Statistical Modeling of Extreme Values*. Springer Series in Statistics. London: Springer.
- Cressie, Noel A. C., and Christopher K. Wikle. 2011. *Statistics for Spatio-Temporal Data*. Hoboken, N.J.: Wiley.
- Downey, Allen B. 2021. *Think Bayes*. "O'Reilly Media, Inc.". <https://alldowney.github.io/ThinkBayes2/>.
- Farnham, David J, James Doss-Gollin, and Upmanu Lall. 2018. “Regional Extreme Precipitation Events: Robust Inference from Credibly Simulated GCM Variables.” *Water Resources Research* 54 (6). <https://doi.org/10.1002/2017wr021318>.
- Friedman, Jerome, Trevor Hastie, and Robert Tibshirani. 2001. *The Elements of Statistical Learning*. Vol. 1. Springer series in statistics Springer, Berlin.
- Gelman, Andrew. 2021. *Regression and Other Stories*. Analytical Methods for Social Research. Cambridge, United Kingdom ; Cambridge University Press.
- Gelman, Andrew, John B Carlin, Hal S Stern, and Donald B Rubin. 2014. *Bayesian Data Analysis*. 3rd ed. Chapman & Hall/CRC Boca Raton, FL, USA.
- Ghil, M, P Yiou, S Hallegatte, B D Malamud, P Naveau, A Soloviev, P Friederichs, et al. 2011. “Extreme Events: Dynamics, Statistics and Prediction.” *Nonlinear Processes in Geophysics* 18 (3): 295–350. <https://doi.org/10/fvzxvv>.
- Helsel, Dennis R., Robert M. Hirsch, Karen R. Ryberg, Stacey A. Archfield, and Edward J. Gilroy. 2020. *Statistical Methods in Water Resources. Techniques and Methods*. U.S. Geological Survey. <https://doi.org/10.3133/tm4A3>.
- Herman, Jon, and Will Usher. 2017. “SALib: An Open-Source Python Library for Sensitivity Analysis.” *Journal of Open Source Software* 2 (9): 97. <https://doi.org/10.21105/joss.00097>.
- Jaynes, Edwin T. 2003. *Probability Theory: The Logic of Science*. New York, NY: Cambridge University Press.
- Kosmyna, Nataliya, Eugene Hauptmann, Ye Tong Yuan, Jessica Situ, Xian-Hao Liao, Ashly Vivian Beresnitzky, Iris Braunstein, and Pattie Maes. 2025. “Your Brain on ChatGPT: Accumulation

- of Cognitive Debt When Using an AI Assistant for Essay Writing Task.” June 10, 2025. <https://doi.org/10.48550/arXiv.2506.08872>.
- Lafferty, David C., and Ryan L. Sriver. 2023. “Downscaling and Bias-Correction Contribute Considerable Uncertainty to Local Climate Projections in CMIP6.” *Npj Climate and Atmospheric Science* 6 (1, 1): 1–13. <https://doi.org/10.1038/s41612-023-00486-0>.
- Lanzante, John R, Keith W Dixon, Mary Jo Nath, Carolyn E Whitlock, and Dennis Adams-Smith. 2018. “Some Pitfalls in Statistical Downscaling of Future Climate.” *Bulletin of the American Meteorological Society* 99 (4): 791–803. <https://doi.org/10.1175/bams-d-17-0046.1>.
- McElreath, Richard. 2020. *Statistical Rethinking: A Bayesian Course with Examples in R and Stan*. Second edition. Texts in Statistical Science Series. Boca Raton ; CRC Press, Taylor & Francis Group.
- Merz, Bruno, Jeroen C J H Aerts, Karsten Arnbjerg-Nielsen, M Baldi, A Becker, A Bichet, Günter Blöschl, et al. 2014. “Floods and Climate: Emerging Perspectives for Flood Risk Assessment and Management.” *Natural Hazards and Earth System Science* 14 (7): 1921–42. <https://doi.org/10/gb9nzm>.
- Mignan, Arnaud. 2024. *Introduction to Catastrophe Risk Modelling: A Physics-based Approach*. Cambridge: Cambridge University Press. <https://doi.org/10.1017/9781009437370>.
- Mudelsee, Manfred. 2020. “Statistical Analysis of Climate Extremes / Manfred Mudelsee.” In *Statistical Analysis of Climate Extremes*. Cambridge, United Kingdom ; Cambridge University Press.
- Naghetini, Mauro, ed. 2017. *Fundamentals of Statistical Hydrology*. Cham: Springer International Publishing. <https://doi.org/10.1007/978-3-319-43561-9>.
- Piironen, Juho, and Aki Vehtari. 2017. “Comparison of Bayesian Predictive Methods for Model Selection.” *Statistics and Computing* 27 (3): 711–35. <https://doi.org/10.1007/s11222-016-9649-y>.
- Pyrz, Michael J. 2024. *Applied Machine Learning in Python: A Hands-on Guide with Code*. https://geostatsguy.github.io/MachineLearningDemos_Book.
- Rackauckas, Christopher, Yingbo Ma, Julius Martensen, Collin Warner, Kirill Zubov, Rohit Sudekhar, Dominic Skinner, Ali Ramadhan, and Alan Edelman. 2020. “Universal Differential Equations for Scientific Machine Learning.” 2020. <https://doi.org/10.48550/ARXIV.2001.04385>.
- Saltelli, Andrea, Marco Ratto, Terry Andres, Francesca Campolongo, Jessica Cariboni, Debora Gatelli, Michaela Saisana, and Stefano Tarantola. 2008. *Global Sensitivity Analysis: The Primer*. John Wiley & Sons, Ltd. <http://onlinelibrary.wiley.com/doi/abs/10.1002/9780470725184.ch1>.
- Thurey, N., B. Holzschuh, P. Holl, G. Kohl, M. Lino, Q. Liu, P. Schnell, and F. Trost. 2024. *Physics-Based Deep Learning*. <https://physicsbaseddeeplearning.org>.

A Software Setup

If you want to run the computational notebooks in this book, or apply a similar workflow, then these instructions are for you.

A.1 Quick start

This section provides step-by-step instructions to get your development environment set up and running.

A.1.1 Installation steps

1. Install **Visual Studio Code** - your code editor
 - There are other good IDEs out there, and you can absolutely use one.
 - VS Code is a good and well-supported starting point
2. Install **Quarto** - for creating documents with code
 - For step 1, choose your operating system
 - For step 2, choose VS Code as your tool
3. Install **Julia** using **JuliaUp** - the programming language
 - Follow the directions on the GitHub page based on your operating system
 - Don't worry about the Continuous Integration (CI) section or anything below it
 - Install Julia 1.11 using `juliaup add 1.11`
 - Set this to be your default version using `juliaup default 1.11`
 - You should get a message that says something like `Configured the default Julia version to be '1.11'`
4. In **VS Code**: Install **extensions** from the Extensions marketplace
 - Install the **Julia** extension (provides syntax highlighting, code completion, and integrated REPL)
 - Install the **Quarto** extension (provides syntax highlighting and preview capabilities for `.qmd` files)
5. Install **GitHub Desktop** - for version control
 - This is optional if you prefer to use git through the command line or another app, but GitHub Desktop is a good default recommendation

A.1.2 Verification

After installation, you should be able to:

- Open VS Code and see the Julia and Quarto extensions listed
- Open a terminal and type `julia` to start the Julia REPL
- Create a new Quarto document (`.qmd` file) in VS Code with syntax highlighting

A.2 Dig deeper

A.2.1 Julia

[Julia](#) is a fast, modern programming language designed for scientific computing. Its syntax closely mirrors mathematical notation, making it intuitive for researchers while delivering performance comparable to C and Fortran.

JuliaUp is the official Julia version manager. It simplifies installation, allows you to maintain multiple Julia versions simultaneously, and keeps your installation current with the latest releases. This is especially useful as the Julia ecosystem evolves rapidly.

See the [Julia page](#) for more.

A.2.2 Quarto

Quarto is a scientific publishing system that enables you to combine code, results, and narrative text in reproducible documents. Think of it as the next generation of R Markdown, but with multi-language support (Julia, Python, R, and more).

This textbook is written in Quarto. Unlike traditional notebooks, Quarto documents are plain text files that render to multiple output formats (HTML, PDF, Word, presentations) while maintaining computational reproducibility.

You can learn more at:

- Official [Tutorial: Hello, Quarto](#) - basic document creation
- Official [Tutorial: Computations](#) - integrating code
- Comprehensive [Quarto documentation](#)

A.2.2.1 Writing with Markdown and math

Quarto uses Markdown syntax with LaTeX math notation. Essential references:

- [Markdown Cheatsheet](#) - basic text formatting
- [LaTeX Cheatsheet](#) - mathematical notation
- [Mathpix Snip](#) - convert equation images to LaTeX code (free tier available)
- [Detexify](#) - draw symbols to find LaTeX commands

A.2.3 Visual Studio Code

Visual Studio Code is a free, open-source code editor developed by Microsoft. Its strength lies in its extensibility—thousands of extensions add language support, debugging capabilities, and productivity tools.

For our workflow, the Julia extension transforms VS Code into a full Julia development environment with syntax highlighting, intelligent code completion, integrated debugging, and a built-in REPL. The Quarto extension provides similar capabilities for computational documents, including live preview and cell execution.

You can learn more at [the official tutorial](#).

A.2.4 Git and GitHub

Git is a distributed version control system that tracks changes in your code over time. GitHub is a cloud-based platform that hosts Git repositories and adds collaboration features like issue tracking, pull requests, and project management.

Version control is essential for reproducible research—it allows you to track changes, collaborate with others, recover from mistakes, and share your work publicly. This textbook itself is maintained on GitHub.

You can learn more at:

- [Git and GitHub for Poets](#) - beginner-friendly video series
- [GitHub Hello World](#) – official docs
- [Version Control](#) - comprehensive guide from MIT’s “Missing Semester”

B Julia Learning Resources

The computational examples in this textbook use the Julia programming language.

B.1 Why Julia?

Julia is a fast, modern, open-source programming language designed for scientific and numerical computing. The language is designed to be fast, dynamic, and easy to use and maintain.

Key advantages for this textbook include:

- **High-Level Syntax:** Julia has a clean and expressive syntax that closely parallels mathematical notation.
- **Performance:** Julia compiles to efficient machine code, achieving speeds comparable to low-level languages like C and Fortran. This solves the “two-language problem,” where you might prototype in a high-level language but need to rewrite for performance.
- **Simplified Dependencies:** Eliminates or reduces the need for dependencies on C and Fortran libraries, which simplifies installation and maintenance.
- **Open-Source and Shareable:** Julia is completely open-source with excellent package management for reproducible research environments.
- **Strong Ecosystem:** Despite being newer, Julia has a rapidly growing ecosystem of high-quality libraries for scientific domains.

While Julia is powerful for computational thinking and research, many ecosystems remain stronger in other languages (like Python’s deep learning and climate data analysis tools), so a well-rounded programmer benefits from learning multiple languages.

You can read more about Julia’s design philosophy:

- [Julia Data Science](#) textbook is didactic and clear
- [Why We Created Julia](#) from the founders
- [Why Julia Manifesto](#) is more comprehensive

B.2 Learning resources

This textbook aims to never reinvent the wheel. There are lots of exceptional resources for learning Julia, or for learning computational concepts with Juila. Here are some favorites:

- MIT’s [Introduction to Computational Thinking](#): Julia-based course covering applied mathematics and computational thinking
- [Julia for Nervous Beginners](#): free course for people hesitant but curious about learning Julia

- [Julia Data Science](#): comprehensive introduction to data science with Julia
- [FastTrack to Julia cheatsheet](#)
- [Comprehensive Julia Tutorials](#): YouTube playlist covering Julia topics
- [Matlab-Python-Julia Cheatsheet](#): helpful if you're experienced in one of these languages

B.2.1 Specialized topics

Here are some additional resources for specific Julia tools and packages developed in this class

- Plotting: [Makie Tutorials](#) and [MakieCon 2023 YouTube Channel](#)
- Statistical Modeling: [Turing.jl tutorials](#) has detailed examples of using Turing for modeling

C Large Language Models (“AI”)

Coding is an integral part of real-world climate-risk analysis, and large language models (LLMs; often referred to as “AI” models) are rapidly changing how some kinds of coding happen. Beyond web-based chatbots, you may have used tools like [GitHub Copilot](#) (free for students and educators) or Claude Code (see free [Deeplearning.AI Course](#)). LLMs use powerful new technologies that can support learning and replace tedious tasks, but they can also threaten your intellectual growth and skill development (Kosmyna et al. 2025; Bastani et al. 2025).

It is clear that there are some tasks that should be delegated to these models and some tasks that must remain human-driven. However, there are tremendous differences of opinion about how most tasks in the middle can or should be allocated. As you wrestle with these questions for yourself, you should explore resources like:

- [AI Snake Oil](#) is a blog that seeks to dispel hype, remove misconceptions, and clarify the limits of AI. The authors are in the Princeton University Department of Computer Science.
- [AI software assistants make the hardest kinds of bugs to spot](#) from Pluralistic is a thoughtful and deep blog post about the perils of (mis)using LLMs for coding.
- [One Useful Thing](#) is a newsletter about AI focused on implications for work and education. The authors’ [prompt library](#) is also a good resource for working with LLMs.
- [Ed Zitron’s Where’s Your Ed At](#) is a newsletter that takes a critical perspective on the business models and hype narratives around AI.